

# Transaction Management

Advanced Database Systems

[ 1 ]



## When *Events* occur...



Hardware and  
Software failure



*Multiple users are  
accessing the DB*



[ 2 ]



# DBMS should ensure that...

## A Database

- ***remains Reliable***
- ***In a Consistent State***



[ 3 ]



# DBMS Functions

- Transaction support
- Concurrency control
- Recovery techniques

Chapter 14: Transactions

Silberschatz, A., Korth, H. & Sudarshan, S. (2010). *Database System Concepts*,  
6th Edition. McGraw-Hill Book Co.

[ 4 ]



# Outline

- Transactions
  - Definition
  - Basic Operations – read and write
  - Properties – A.C.I.D.
  - Transaction Operations
  - States and Outcomes
  - Isolation Levels
- Transaction Definition in SQL

[ 5 ]



# Transaction

- A logical unit of processing
    - Entire program
    - Part of a program
    - A single command
  - Carried out by
    - A single user
    - An application program
  - To access the database
    - Write: Insert, Update, Delete
    - Read: Retrieve
- A sequence of one or more SQL operations treated as a unit

[ 6 ]



# Transaction Boundaries

- Begin transaction
- End transaction

$t_1$ : Withdrawal transaction



Time	T <sub>1</sub>	bal <sub>x</sub>
t <sub>1</sub>	begin_transaction	100
t <sub>2</sub>	read(bal <sub>x</sub> )	100
t <sub>3</sub>	bal <sub>x</sub> =bal <sub>x</sub> - 10	100
t <sub>4</sub>	write(bal <sub>x</sub> )	90
t <sub>5</sub>	end_transaction	90
t <sub>6</sub>	commit	90

[ 7 ]



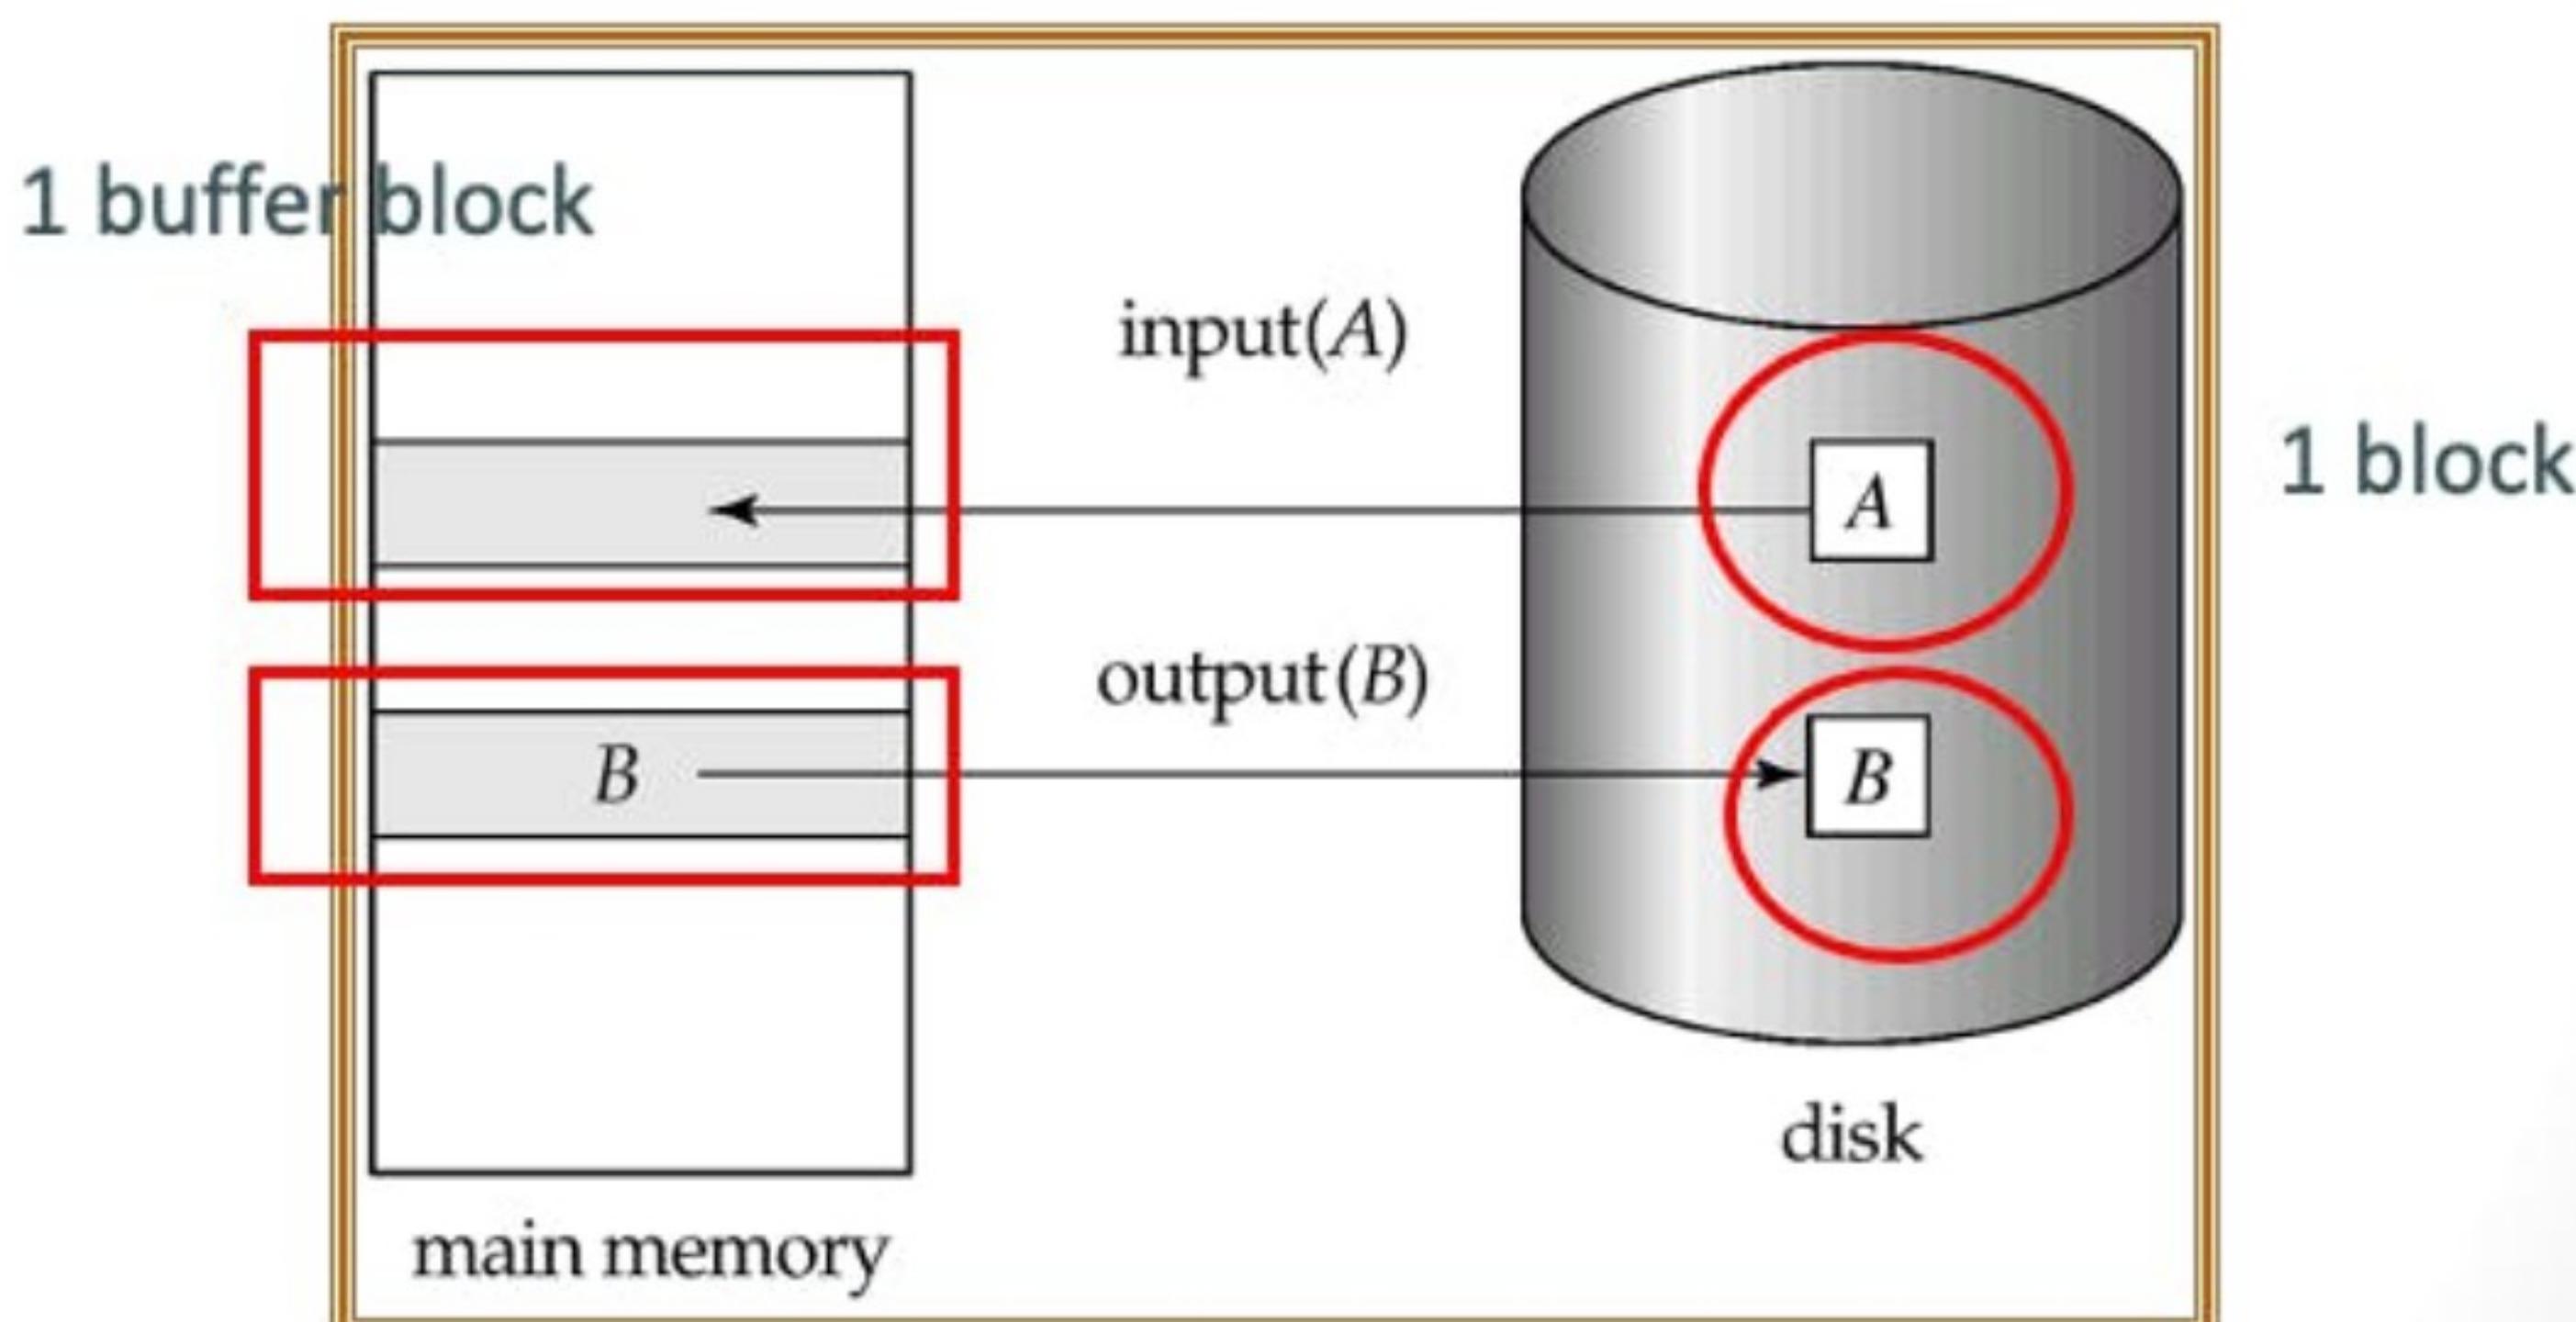
# Simple Model of a Database

- **Database**
  - A collection of named data items
  - Permanently resides on a disk, but some portion of it is temporarily residing in memory
- **Granularity of data (what is read or written)**
  - Field (Column)
  - Record (Tuple)
  - Whole disk block

[ 8 ]



# Block Storage Operations



[ 9 ]



# Basic Transaction Operations

- **read(X)**

- Reads a database item named X into a program variable (local buffer)

- **write(X)**

- Writes the value of program variable (local buffer) into the database item named X

[ 10 ]

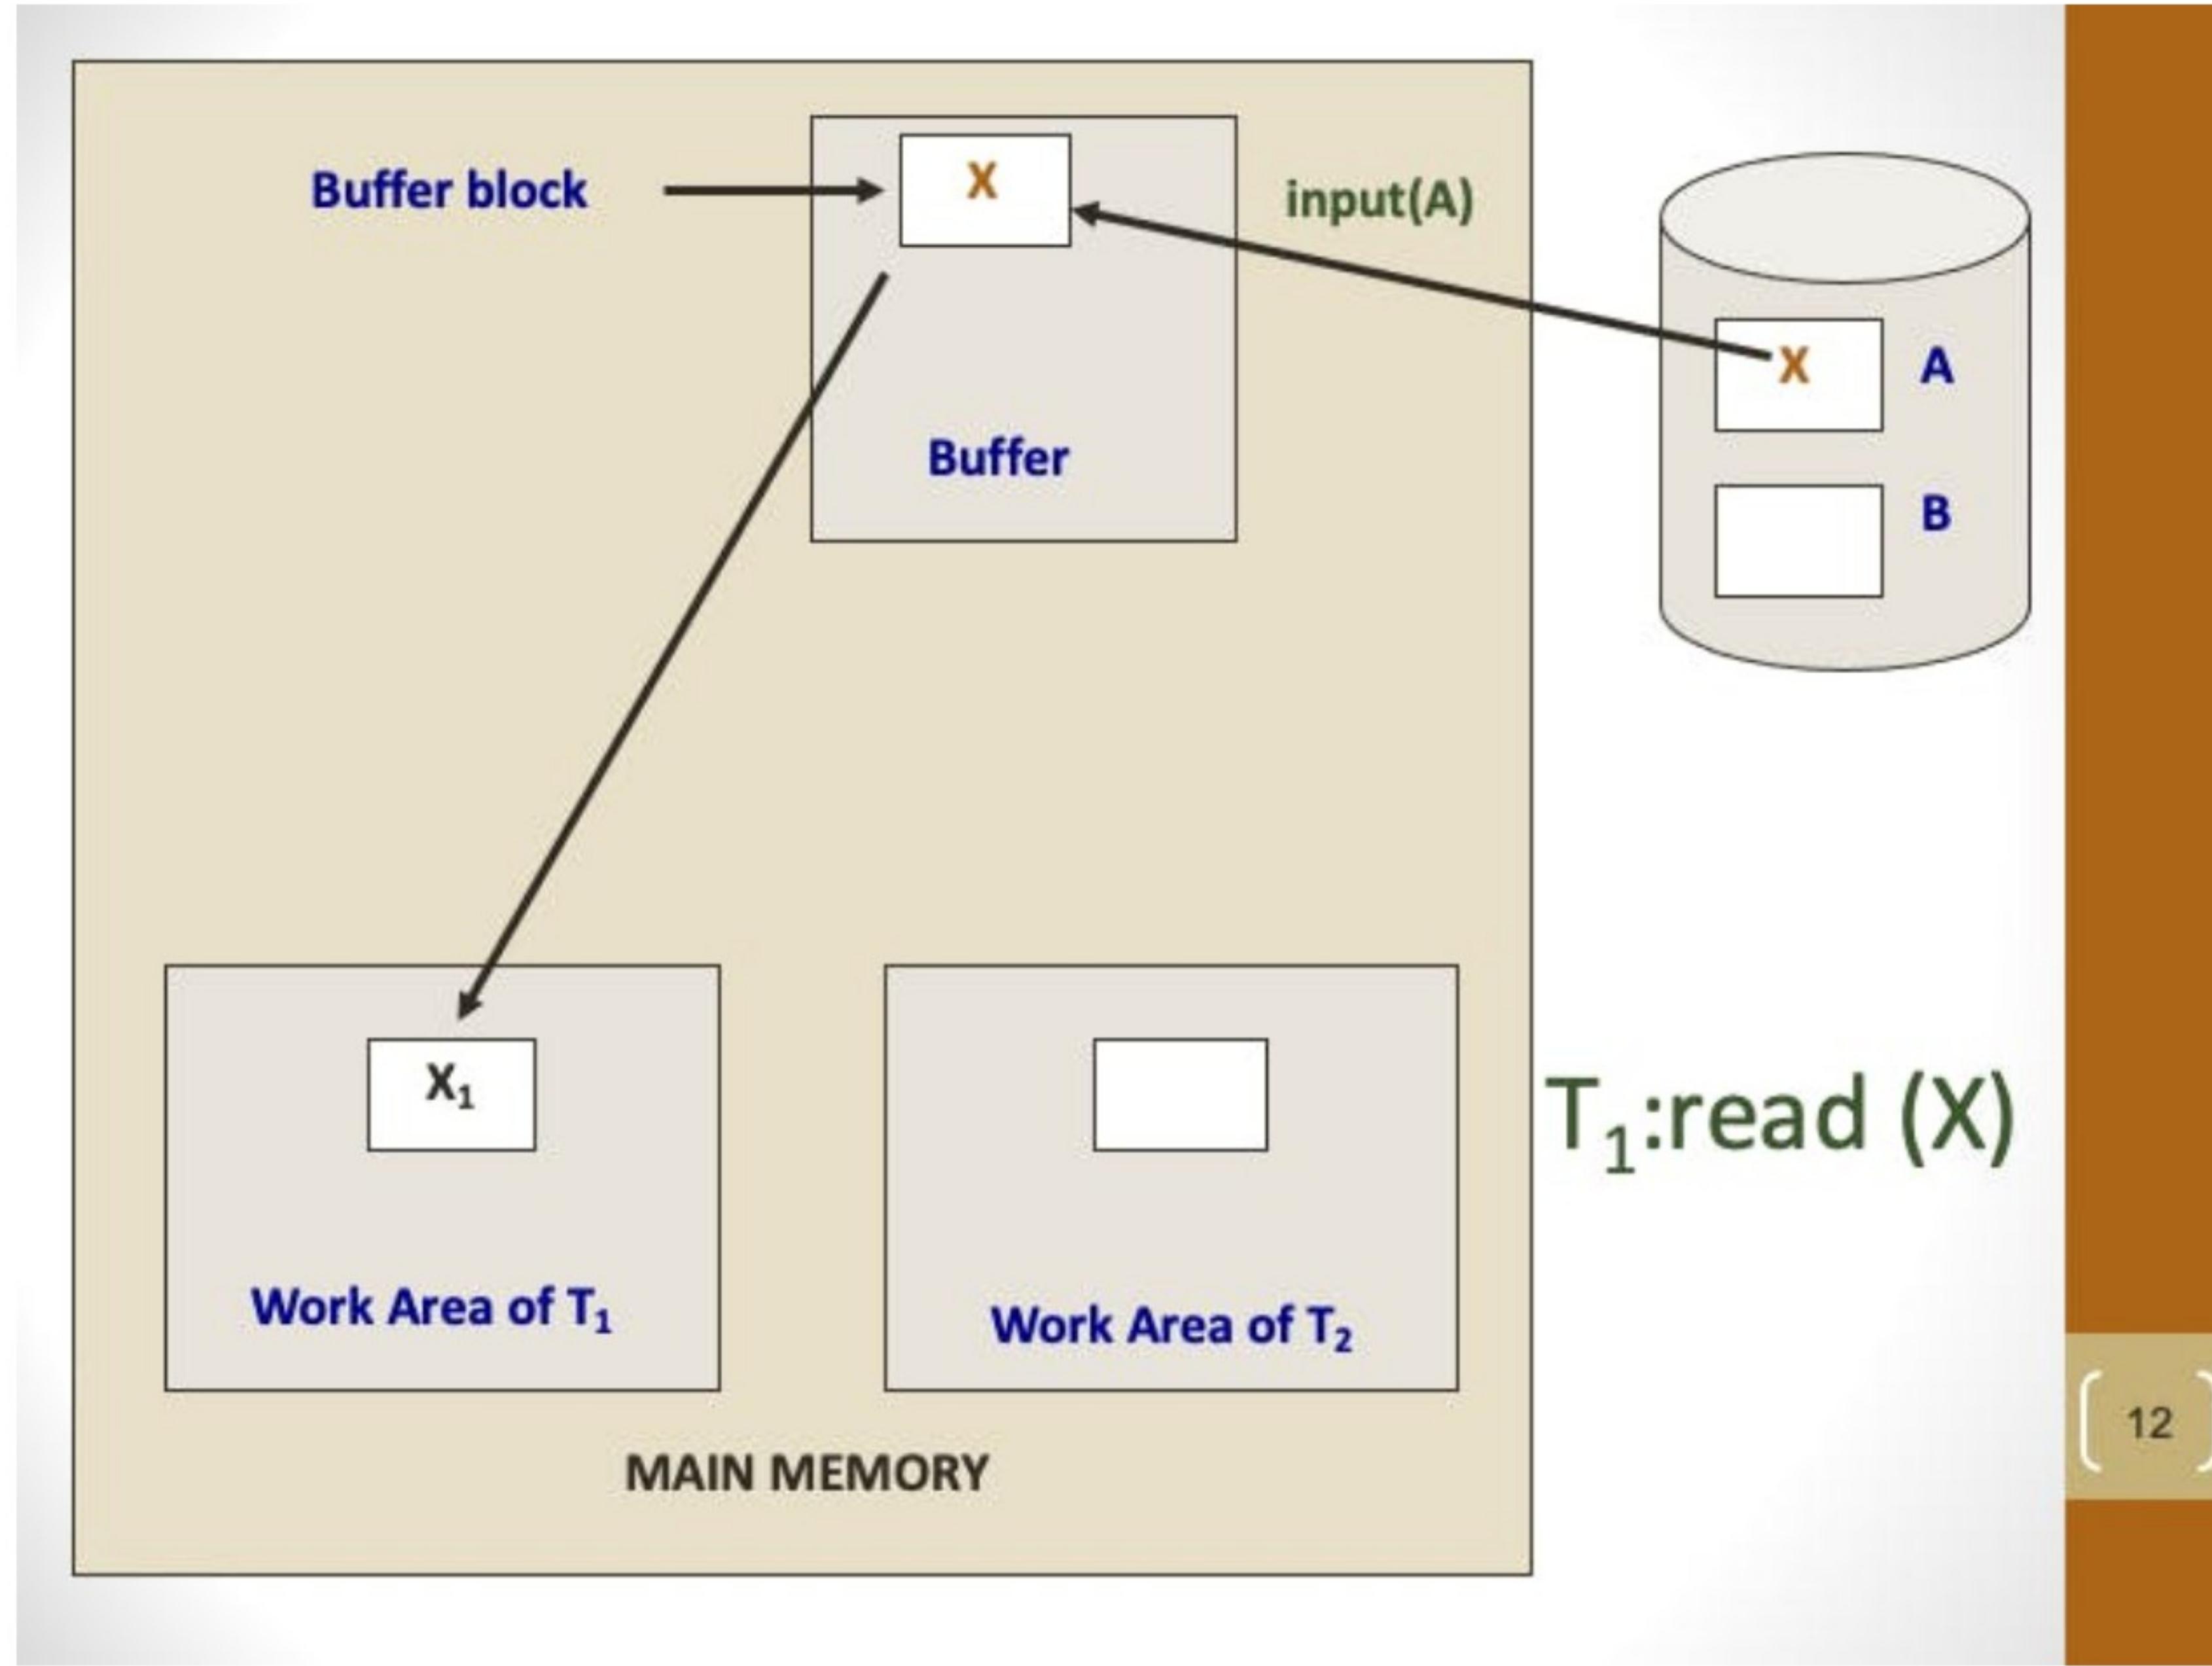


# Steps in Read(X)

1. Find the address of the disk block that contains item X.
2. Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer).
3. Copy item X from the buffer to the program variable.

[ 11 ]



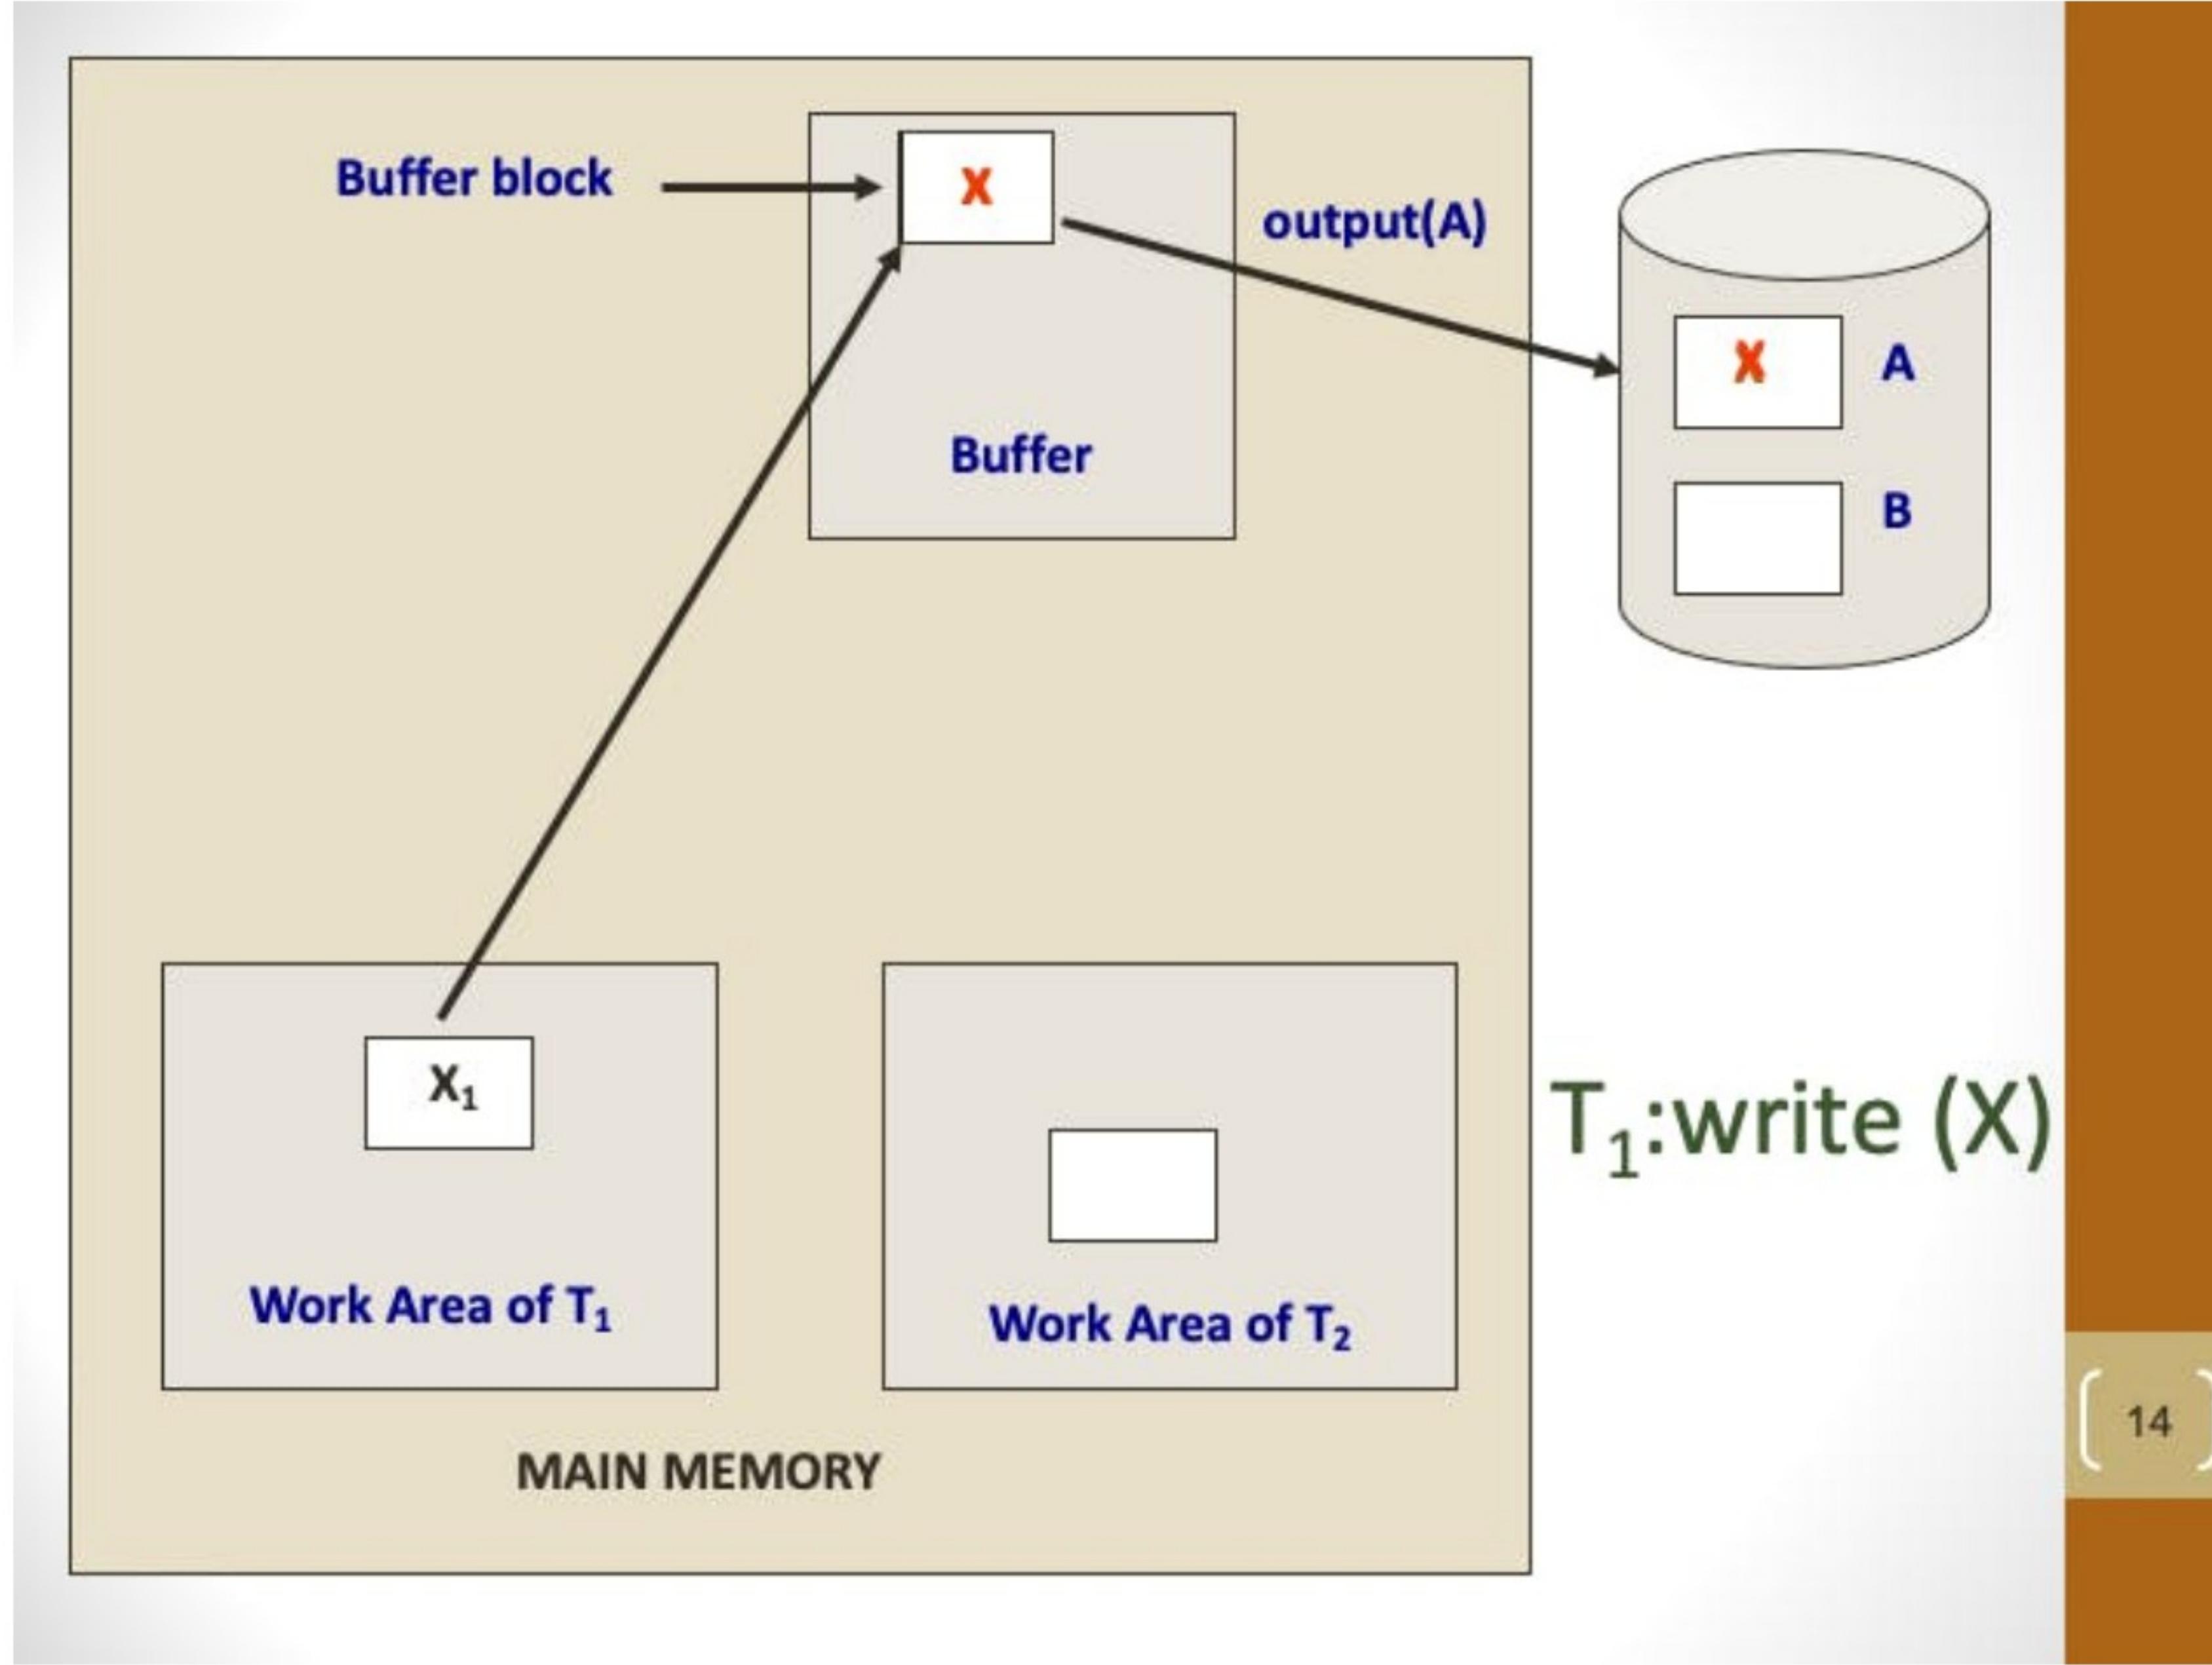


# Steps in Write(X)

1. Find the address of the disk block that contains item X.
2. Copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer).
3. Copy item X from the program variable into its correct location in the buffer.
4. Store the updated block from the buffer back to disk (either immediately or at some later point in time).

[ 13 ]





# Transaction: Examples

Given the following database relations:

EMPLOYEE (Eno, Name, Address, Sex,  
CivilStatus, Bdate, Position,  
Salary, SuperEno, Dno)

PROJECT (Pno, Pname, Plocation, Dno)

WORKS\_ON (Eno, Pno, Hours)

[ 15 ]



# Transaction: Example (1)

Update the salary (add 10% to the current salary) of a particular employee with `Eno = x`.

```
EMPLOYEE (Eno, Name, Address, Sex,  
          CivilStatus, Bdate, Position,  
          Salary, SuperEno, Dno)
```

Transaction consists of 2 dbop and 1 non-dbop.

Read( <code>Eno = x, salary</code> )	dbop
<code>salary = salary * 1.1</code>	
Write( <code>Eno = x, salary</code> )	dbop

[ 16 ]

dbop means “DB Operation”



# Transaction: Example (2)

Delete the employee with Eno=x.

Reassign the projects of the employee with Eno=x to new\_eno.

```
EMPLOYEE (Eno, Name, Address, Sex,  
          CivilStatus, Bdate, Position,  
          Salary, SuperEno, Dno)  
WORKS_ON (Eno, Pno, Hours)
```

[ 17 ]



# Transaction: Example (2)

```
Delete (Eno = x)                                dbop
for all Works_On records
begin
    read (Pno = pno, Eno)                      dbop
    if (Eno = x) then
        begin
            Eno = new_eno
            write (Pno = pno, Eno)      dbop
        end
    end
Write(Eno = x, salary)                           dbop
```

To prevent *Referential Integrity Constraint* violation error,  
**Delete (Eno = x)** should be performed last.



# Transaction Properties (ACID)

- **Atomicity**

- “All or nothing”
- A transaction is an **indivisible unit** that is either performed in its entirety or it is not performed at all.

- **Consistency**

- A transaction must transform the database from one consistent state to another consistent state.
- The database must reflect the real-word at all times.



# Transaction Properties (ACID)

- **Isolation**
  - Transactions execute **independently** of one another.
  - The partial effects of incomplete transactions should not be visible to other transactions.
- **Durability**
  - The effects of a successfully completed (committed) transactions are **permanently** recorded in the database.
  - They must not be lost because of a subsequent failure.

[ 20 ]



# Transaction Properties (ACID)

## An Example

Fund Transfer: Transfer \$50 from Account A to Account B.

Pre-State:	A = \$1000	B = \$2000
Post-State:	A = \$950	B = \$2050

```
read(A);
A := A - 50;
write(A);
read(B);
B := B + 50;
write(B).
```

### CONSISTENCY

- The sum of A and B should be unchanged by the execution of the transaction.
- Ensuring consistency for an individual transaction is the responsibility of the application programmer.



# Transaction Properties (ACID)

## An Example

A = \$1000

B = \$2000

```
read(A);
A := A - 50;
write(A);
read(B);
B := B + 50;
FAIL!
write(B).
```

### ATOMICITY

- The DBMS keeps track (on disk) of the old values of any data on which a transaction performs a write.
- If a transaction does not complete its execution, **the DBMS restores the old values** to make it appear as though the transaction never executed.
- Ensuring atomicity is the **responsibility of the transaction-management component** of the DBMS.



# Transaction Properties (ACID)

## An Example

A = \$1000

B = \$2000

```
read(A);
A := A - 50;
write(A);
read(B);
B := B + 50;
write(B).
```

### DURABILITY

- Updates carried out by the transaction have been written to disk before the transaction completes.
- Information about updates carried out by the transaction and written to disk is sufficient to enable the database to reconstruct the updates when the DBMS is restarted after failure.
- Ensuring durability is the responsibility of the recovery-management component of the DBMS.



# Transaction Properties (ACID)

## An Example

A = \$1000

B = \$2000

```
T1: read(A);
    A := A - 50;
    write(A);
    read(B);
    B := B + 50;
T2: read(A);
    read(B);
    sum = A + B;
    write(B).
```

### ISOLATION

- Execute transactions **serially**; however performance is affected;
- Other solutions have been developed to **allow multiple transactions to execute concurrently**.
- Ensuring isolation is the responsibility of the **concurrency-control component** of the DBMS.



# Transactions in DBMS

- The DBMS does not know which updates are grouped together to form a single transaction.
- The user should indicate the boundaries of the transaction through DML commands (delimiters) provided for in SQL2:
  - BEGIN TRANSACTION
  - COMMIT
- If delimiters are not used, the entire program is considered as a single transaction.

[ 25 ]



# Transaction Operations

T	T <sub>1</sub>	bal <sub>x</sub>
t <sub>1</sub>	begin _transaction	100
t <sub>2</sub>	read(bal <sub>x</sub> )	100
t <sub>3</sub>	bal <sub>x</sub> =bal <sub>x</sub> - 10	100
t <sub>4</sub>	write(bal <sub>x</sub> )	90
t <sub>5</sub>	end_transaction	90
t <sub>6</sub>	commit	90

- BEGIN\_TRANSACTION
  - Marks the beginning of transaction execution
- READ or WRITE
  - Specify the read or write operations on database items that are executed as part of a transaction



# Transaction Operations

- **END\_TRANSACTION**

- Marks the end limit of transaction execution
- Specifies that READ and WRITE transaction operations have ended
- At this point it may be necessary to check whether :
  - The changes introduced by the transaction can be permanently applied (committed) to the database; or
  - The transaction has to be aborted because it violates concurrency control for some other reason.



# Transaction Operations

- **COMMIT**
  - Signals a successful end of the transaction
  - Any changes (updates) executed by the transaction can be safely committed to the database and will not be undone.
- **ROLLBACK**
  - Signals that the transaction has ended unsuccessfully
  - Any changes or effects that the transaction may have applied to the database must be undone.

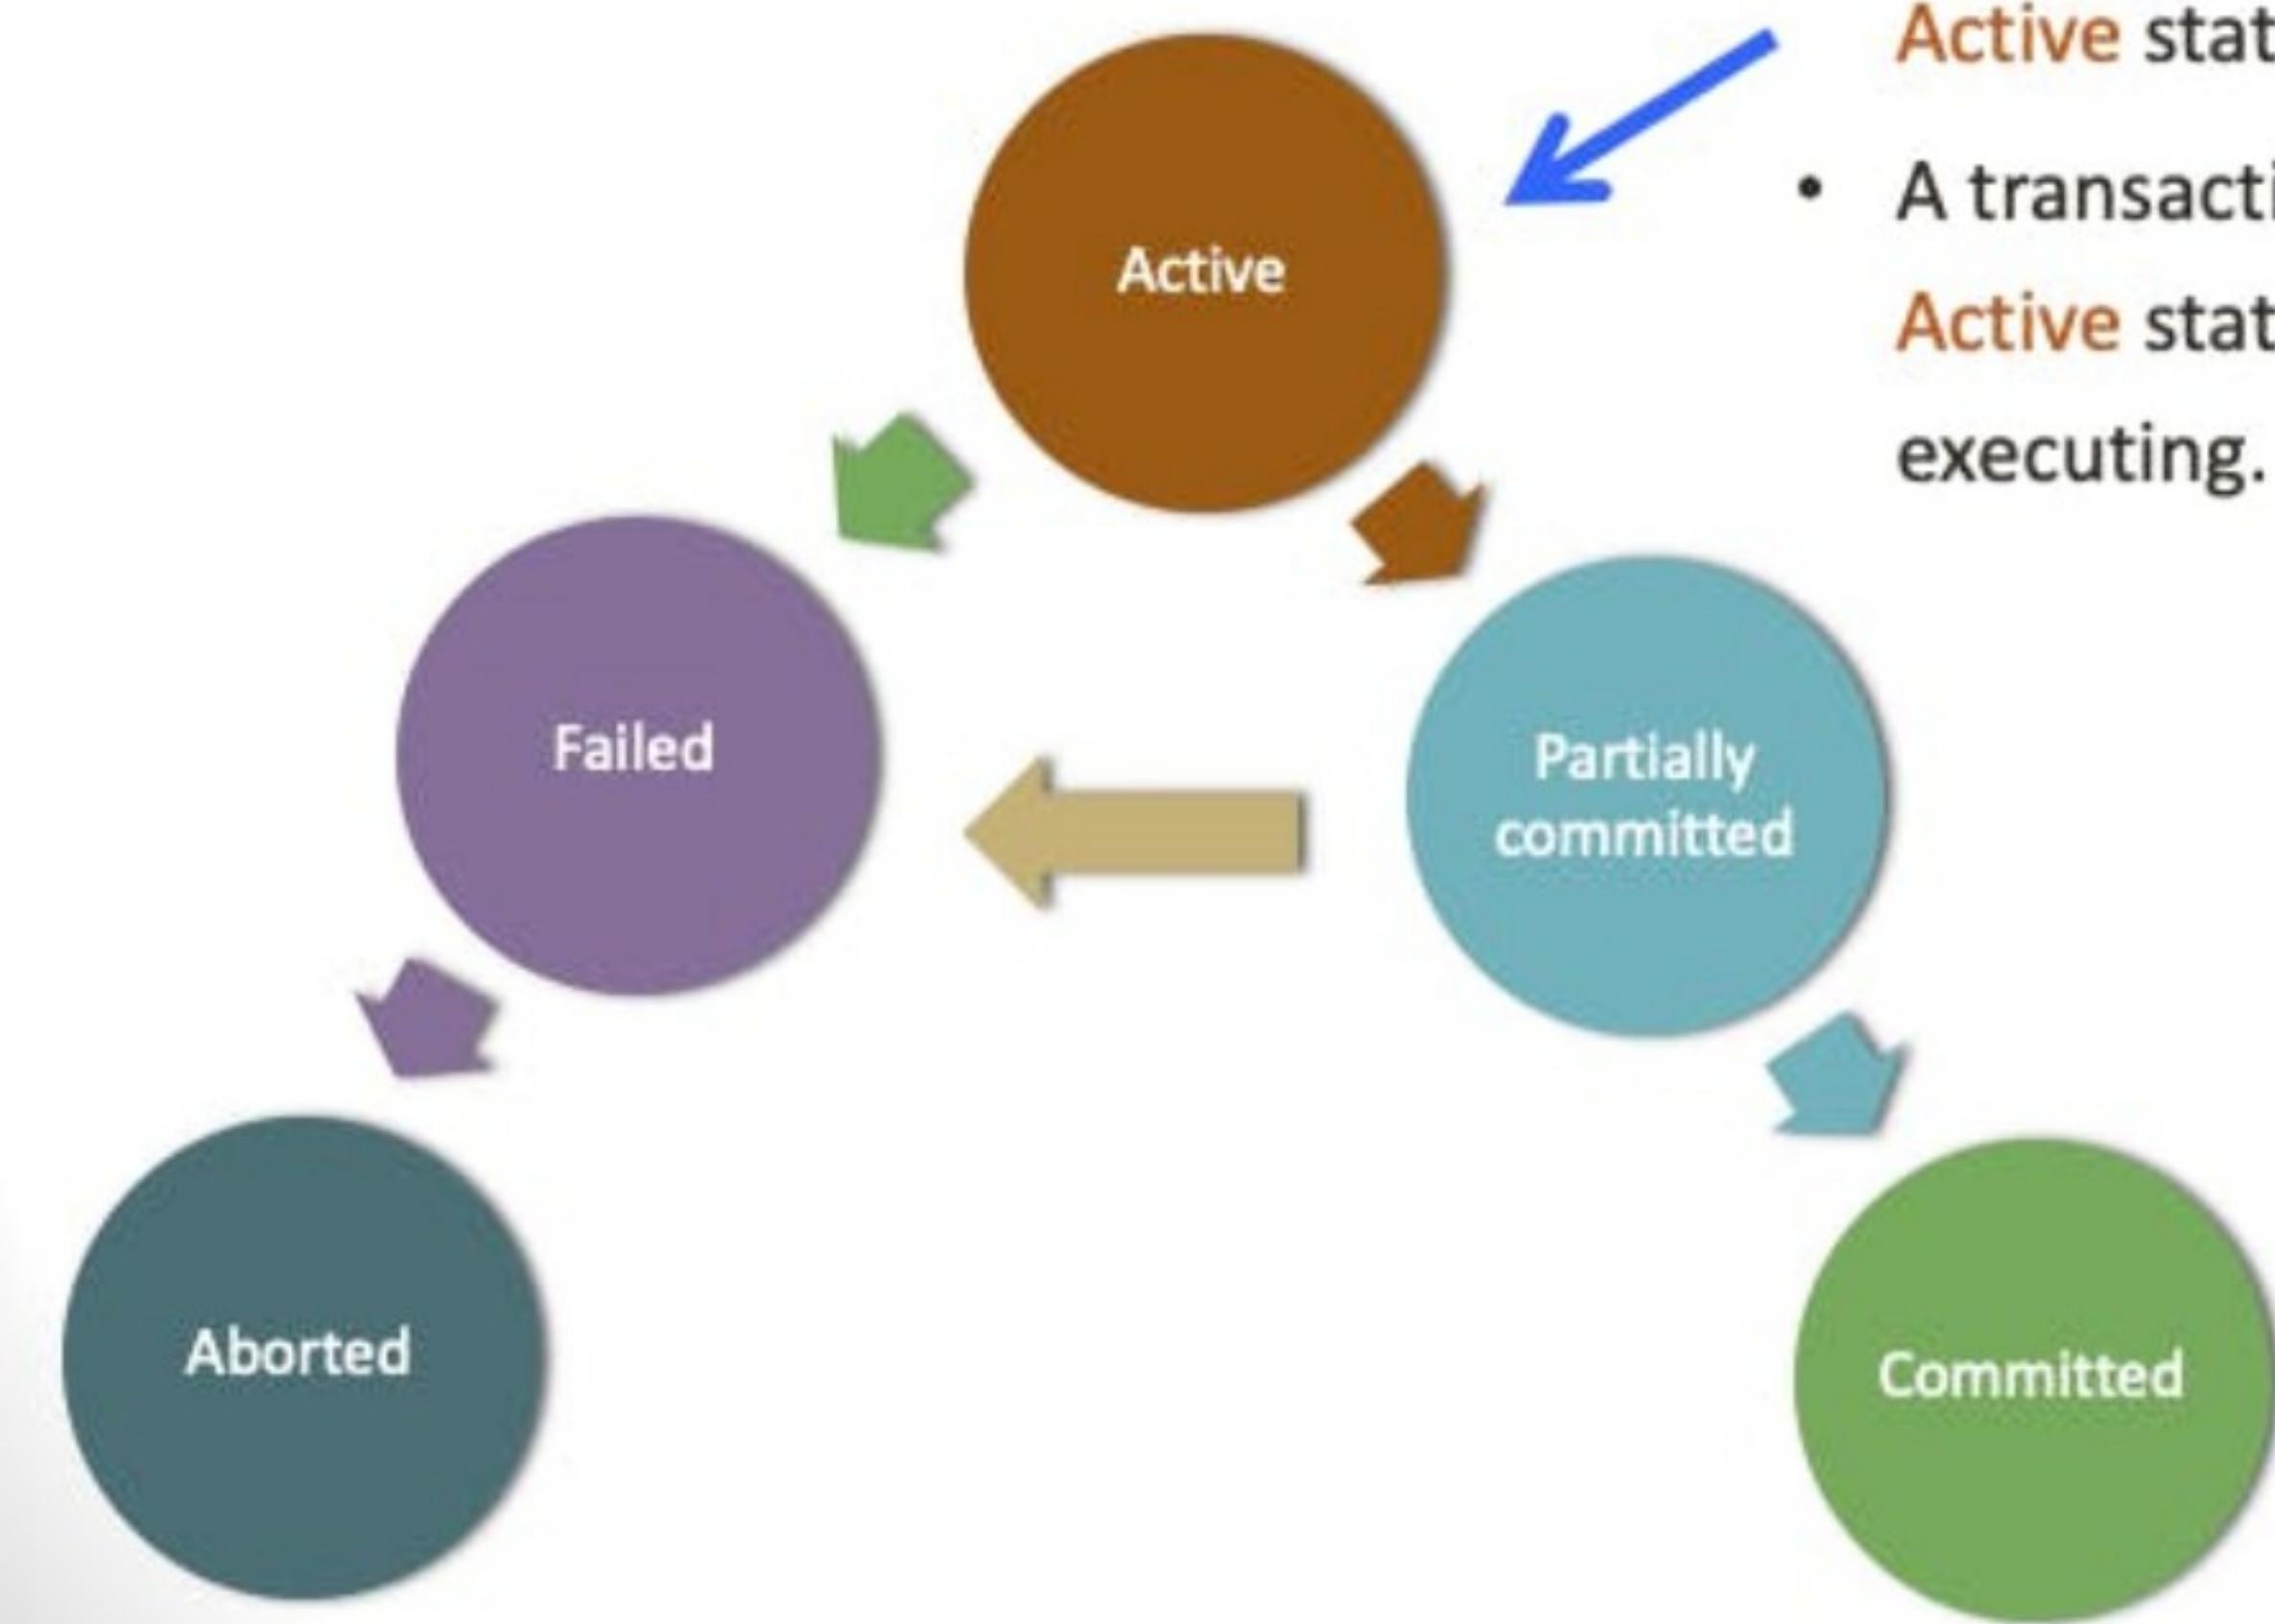


# Transaction Operations

- **UNDO**
  - Similar to rollback except that it applies to a single operation rather than a whole transaction.
- **REDO**
  - This specifies that certain transaction operations must be redone to ensure that all the operations of a committed transaction have been applied successfully to the database.



# Transaction States



- A transaction starts in the **Active** state.
- A transaction stays in the **Active** state while it is executing.

[ 30 ]



# Transaction States: Example

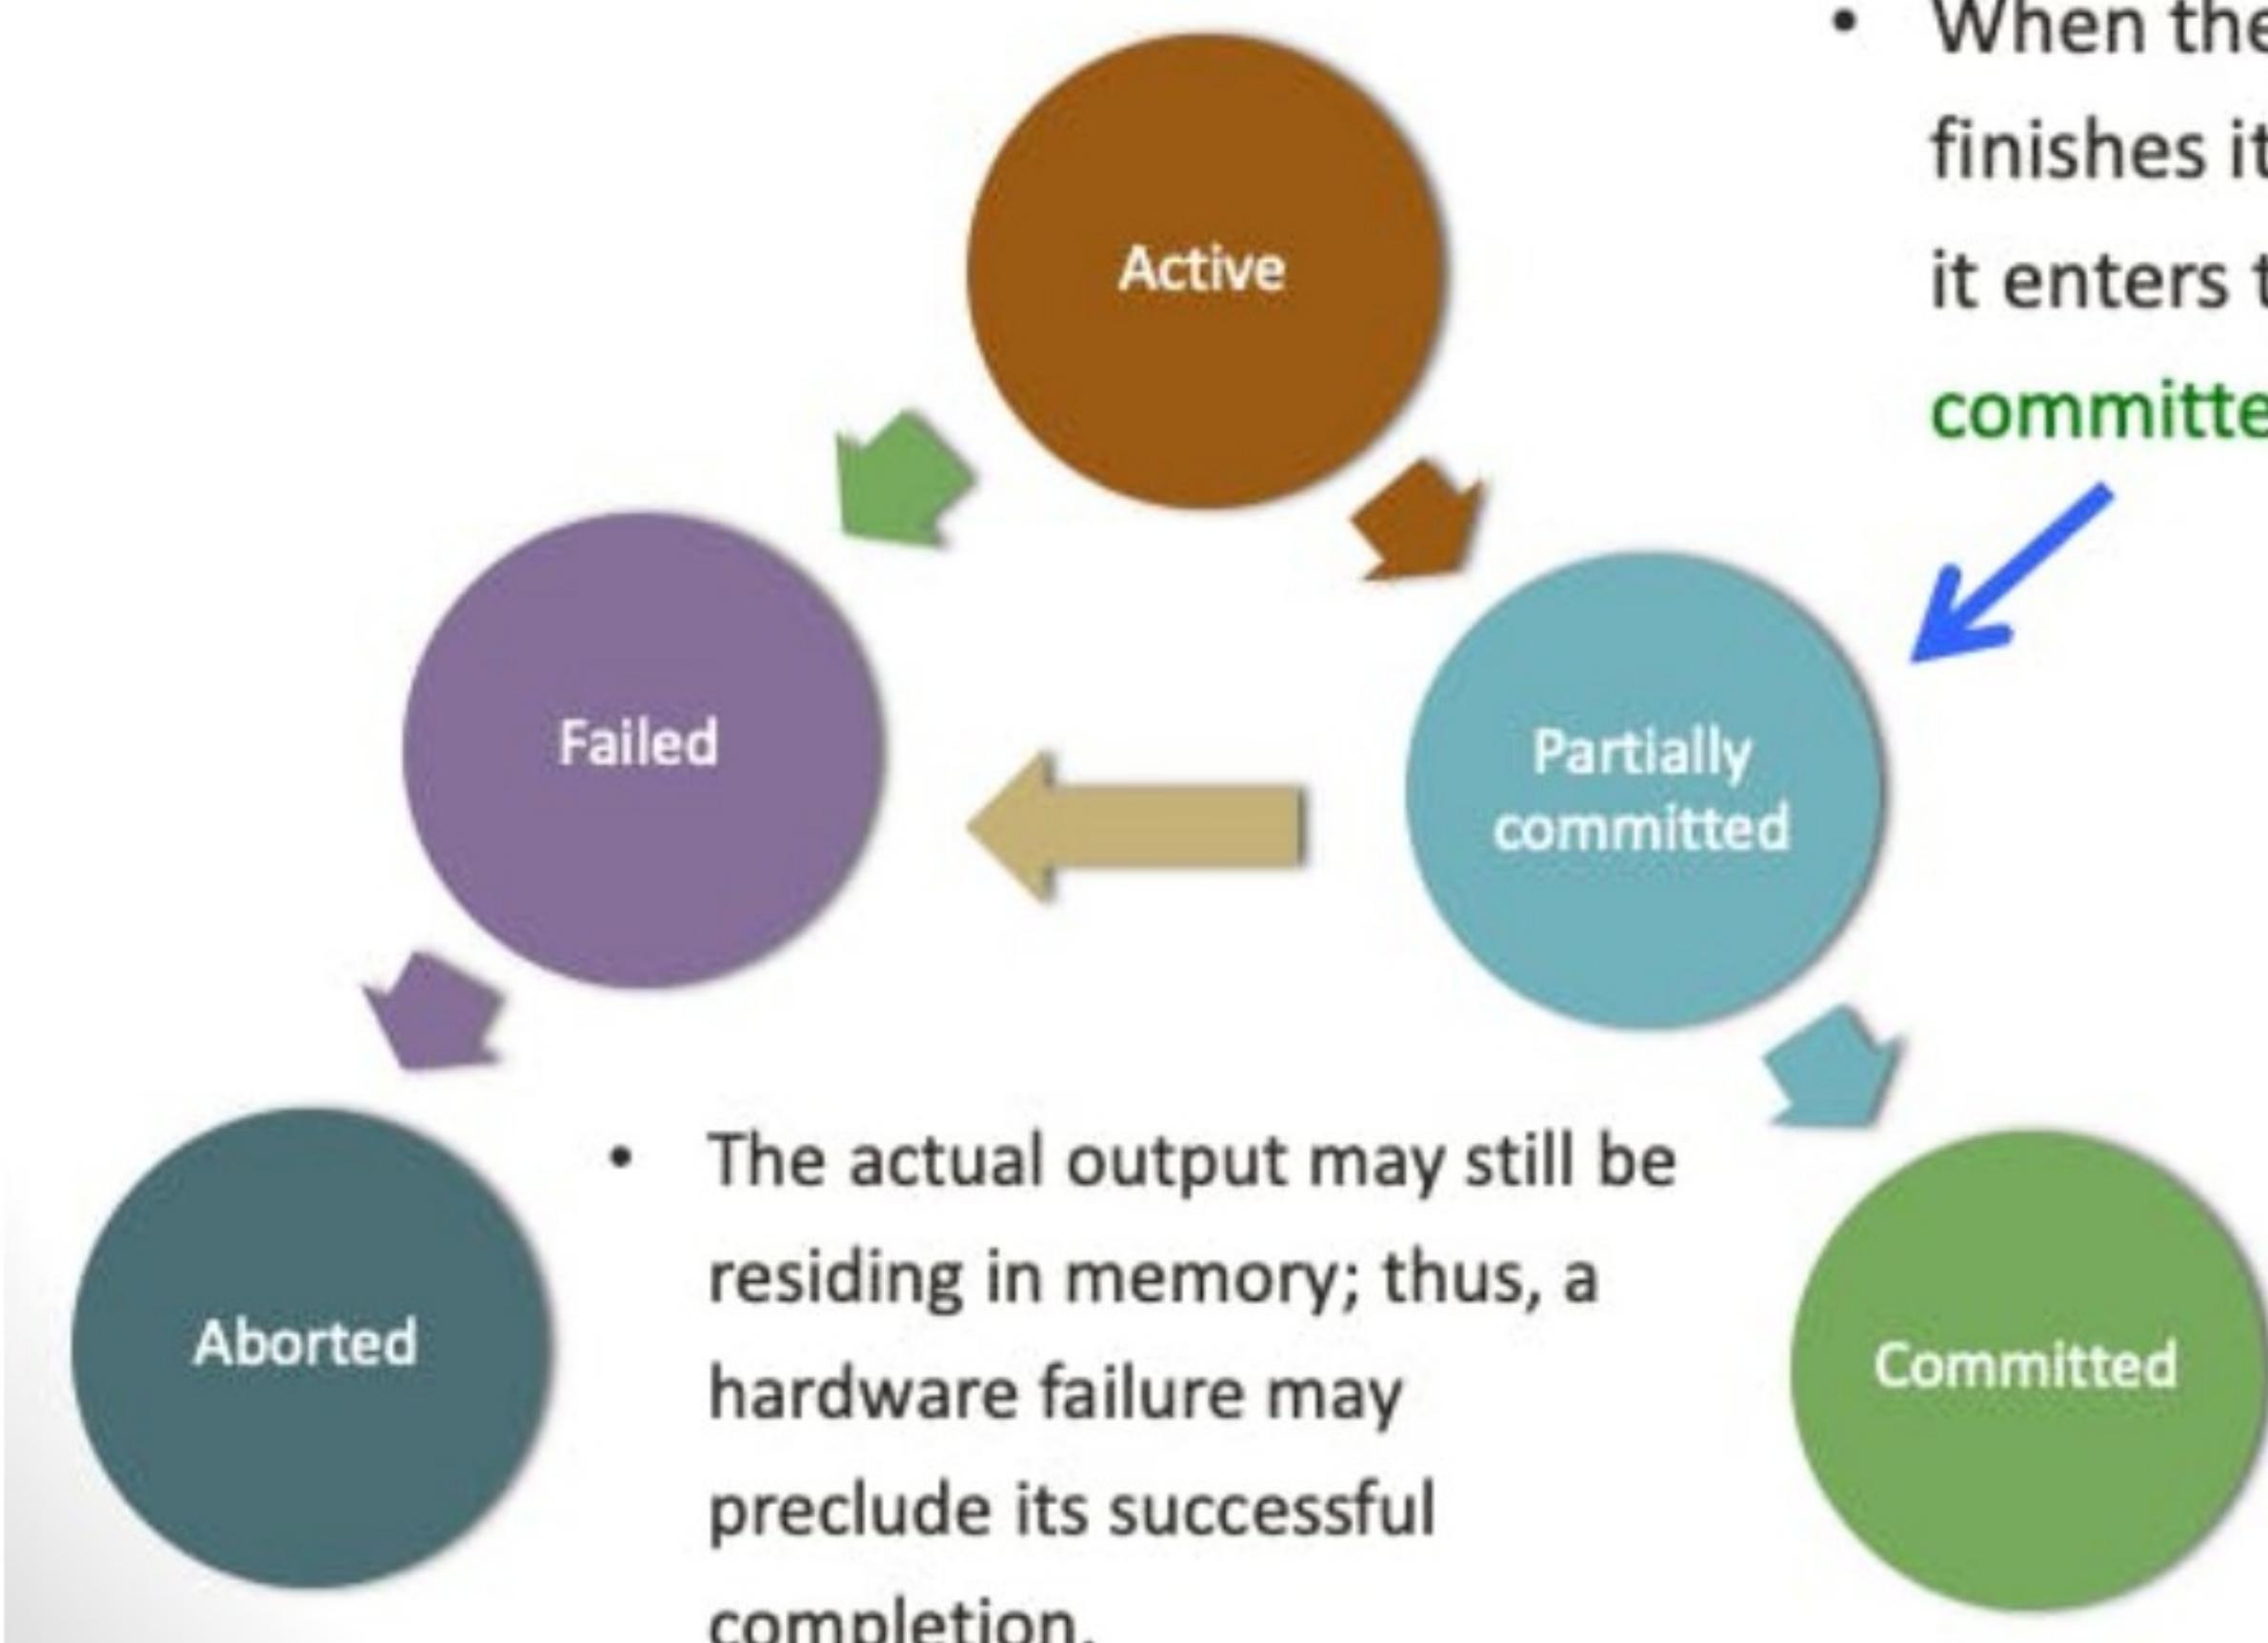
$T_1$  : Withdrawal transaction

Time	$T_1$	$bal_x$
$t_1$	begin_transaction	100
$t_2$	read( $bal_x$ )	100
$t_3$	$bal_x = bal_x - 10$	100
$t_4$	write( $bal_x$ )	90
$t_5$	end_transaction	90
$t_6$	commit	90

Active



# Transaction States



[ 32 ]



# Transaction States: Example

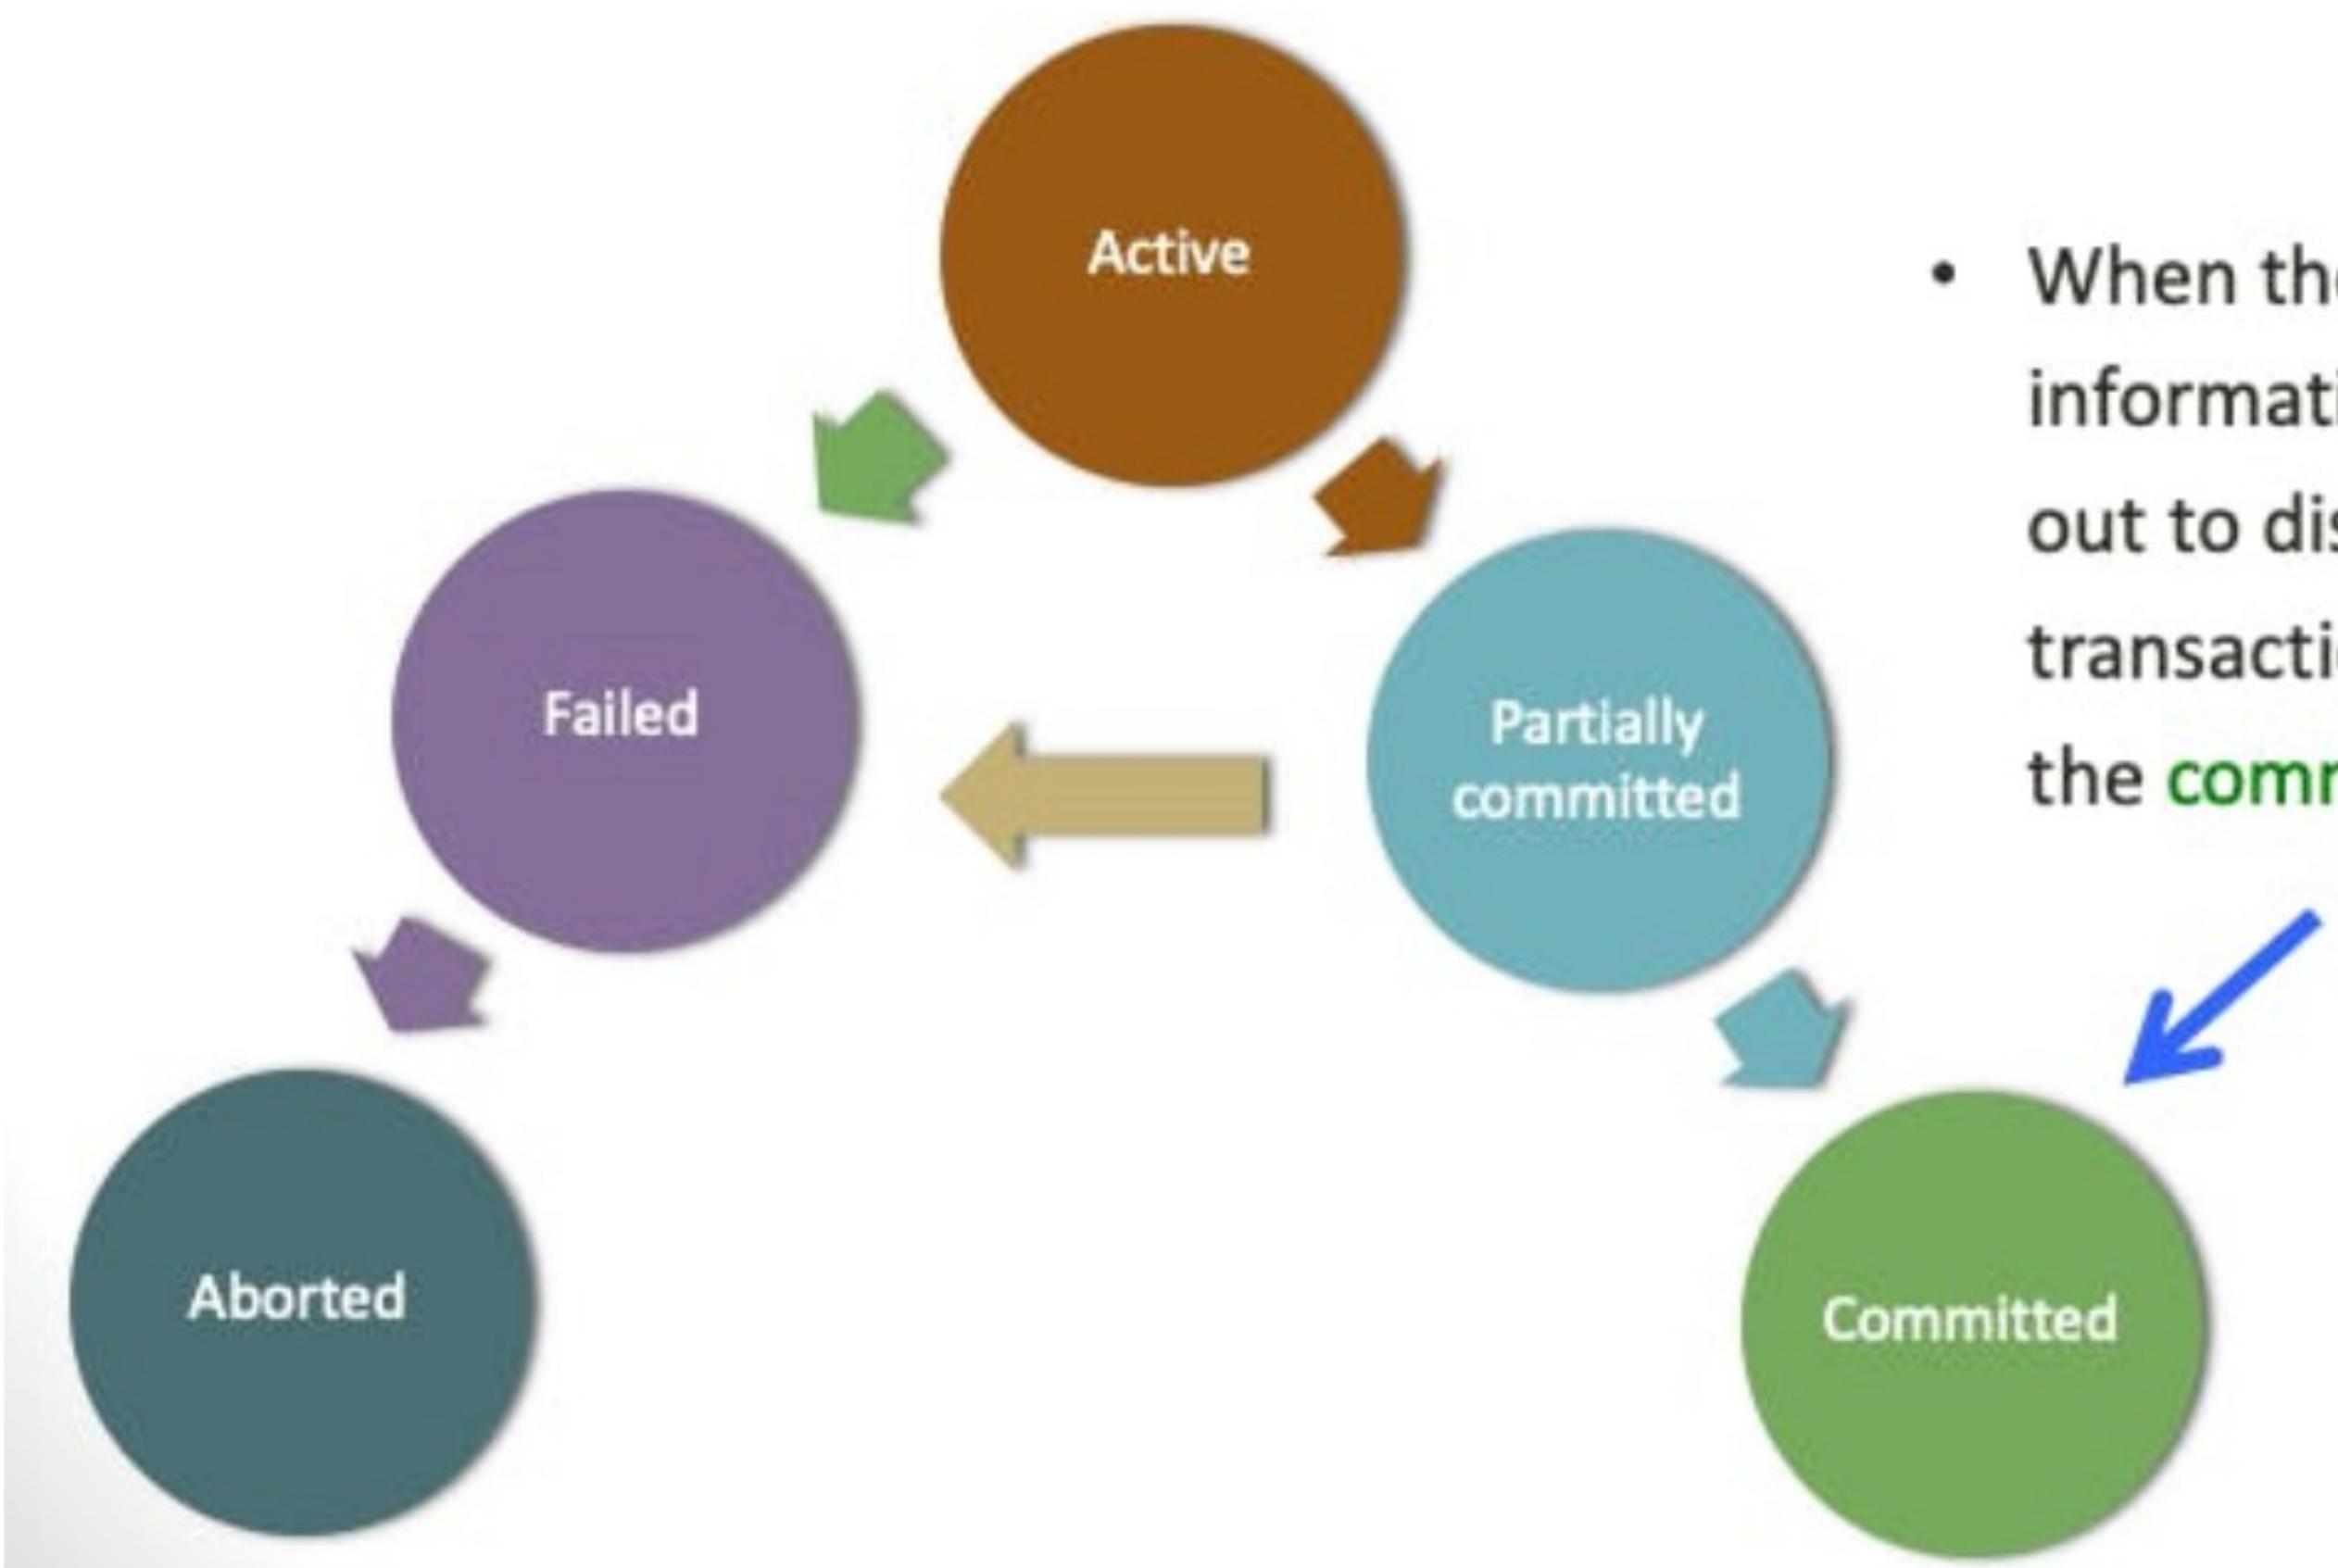
$T_1$  : Withdrawal transaction

Partially completed

Time	$T_1$	$bal_x$
$t_1$	begin _transaction	100
$t_2$	read( $bal_x$ )	100
$t_3$	$bal_x=bal_x - 10$	100
$t_4$	write( $bal_x$ )	90
$t_5$	end _transaction	90
$t_6$	commit	90



# Transaction States



- When the last of the information is written out to disk, the transaction enters the **committed** state.

[ 34 ]



# Transaction States: Example

$T_1$  : Withdrawal transaction

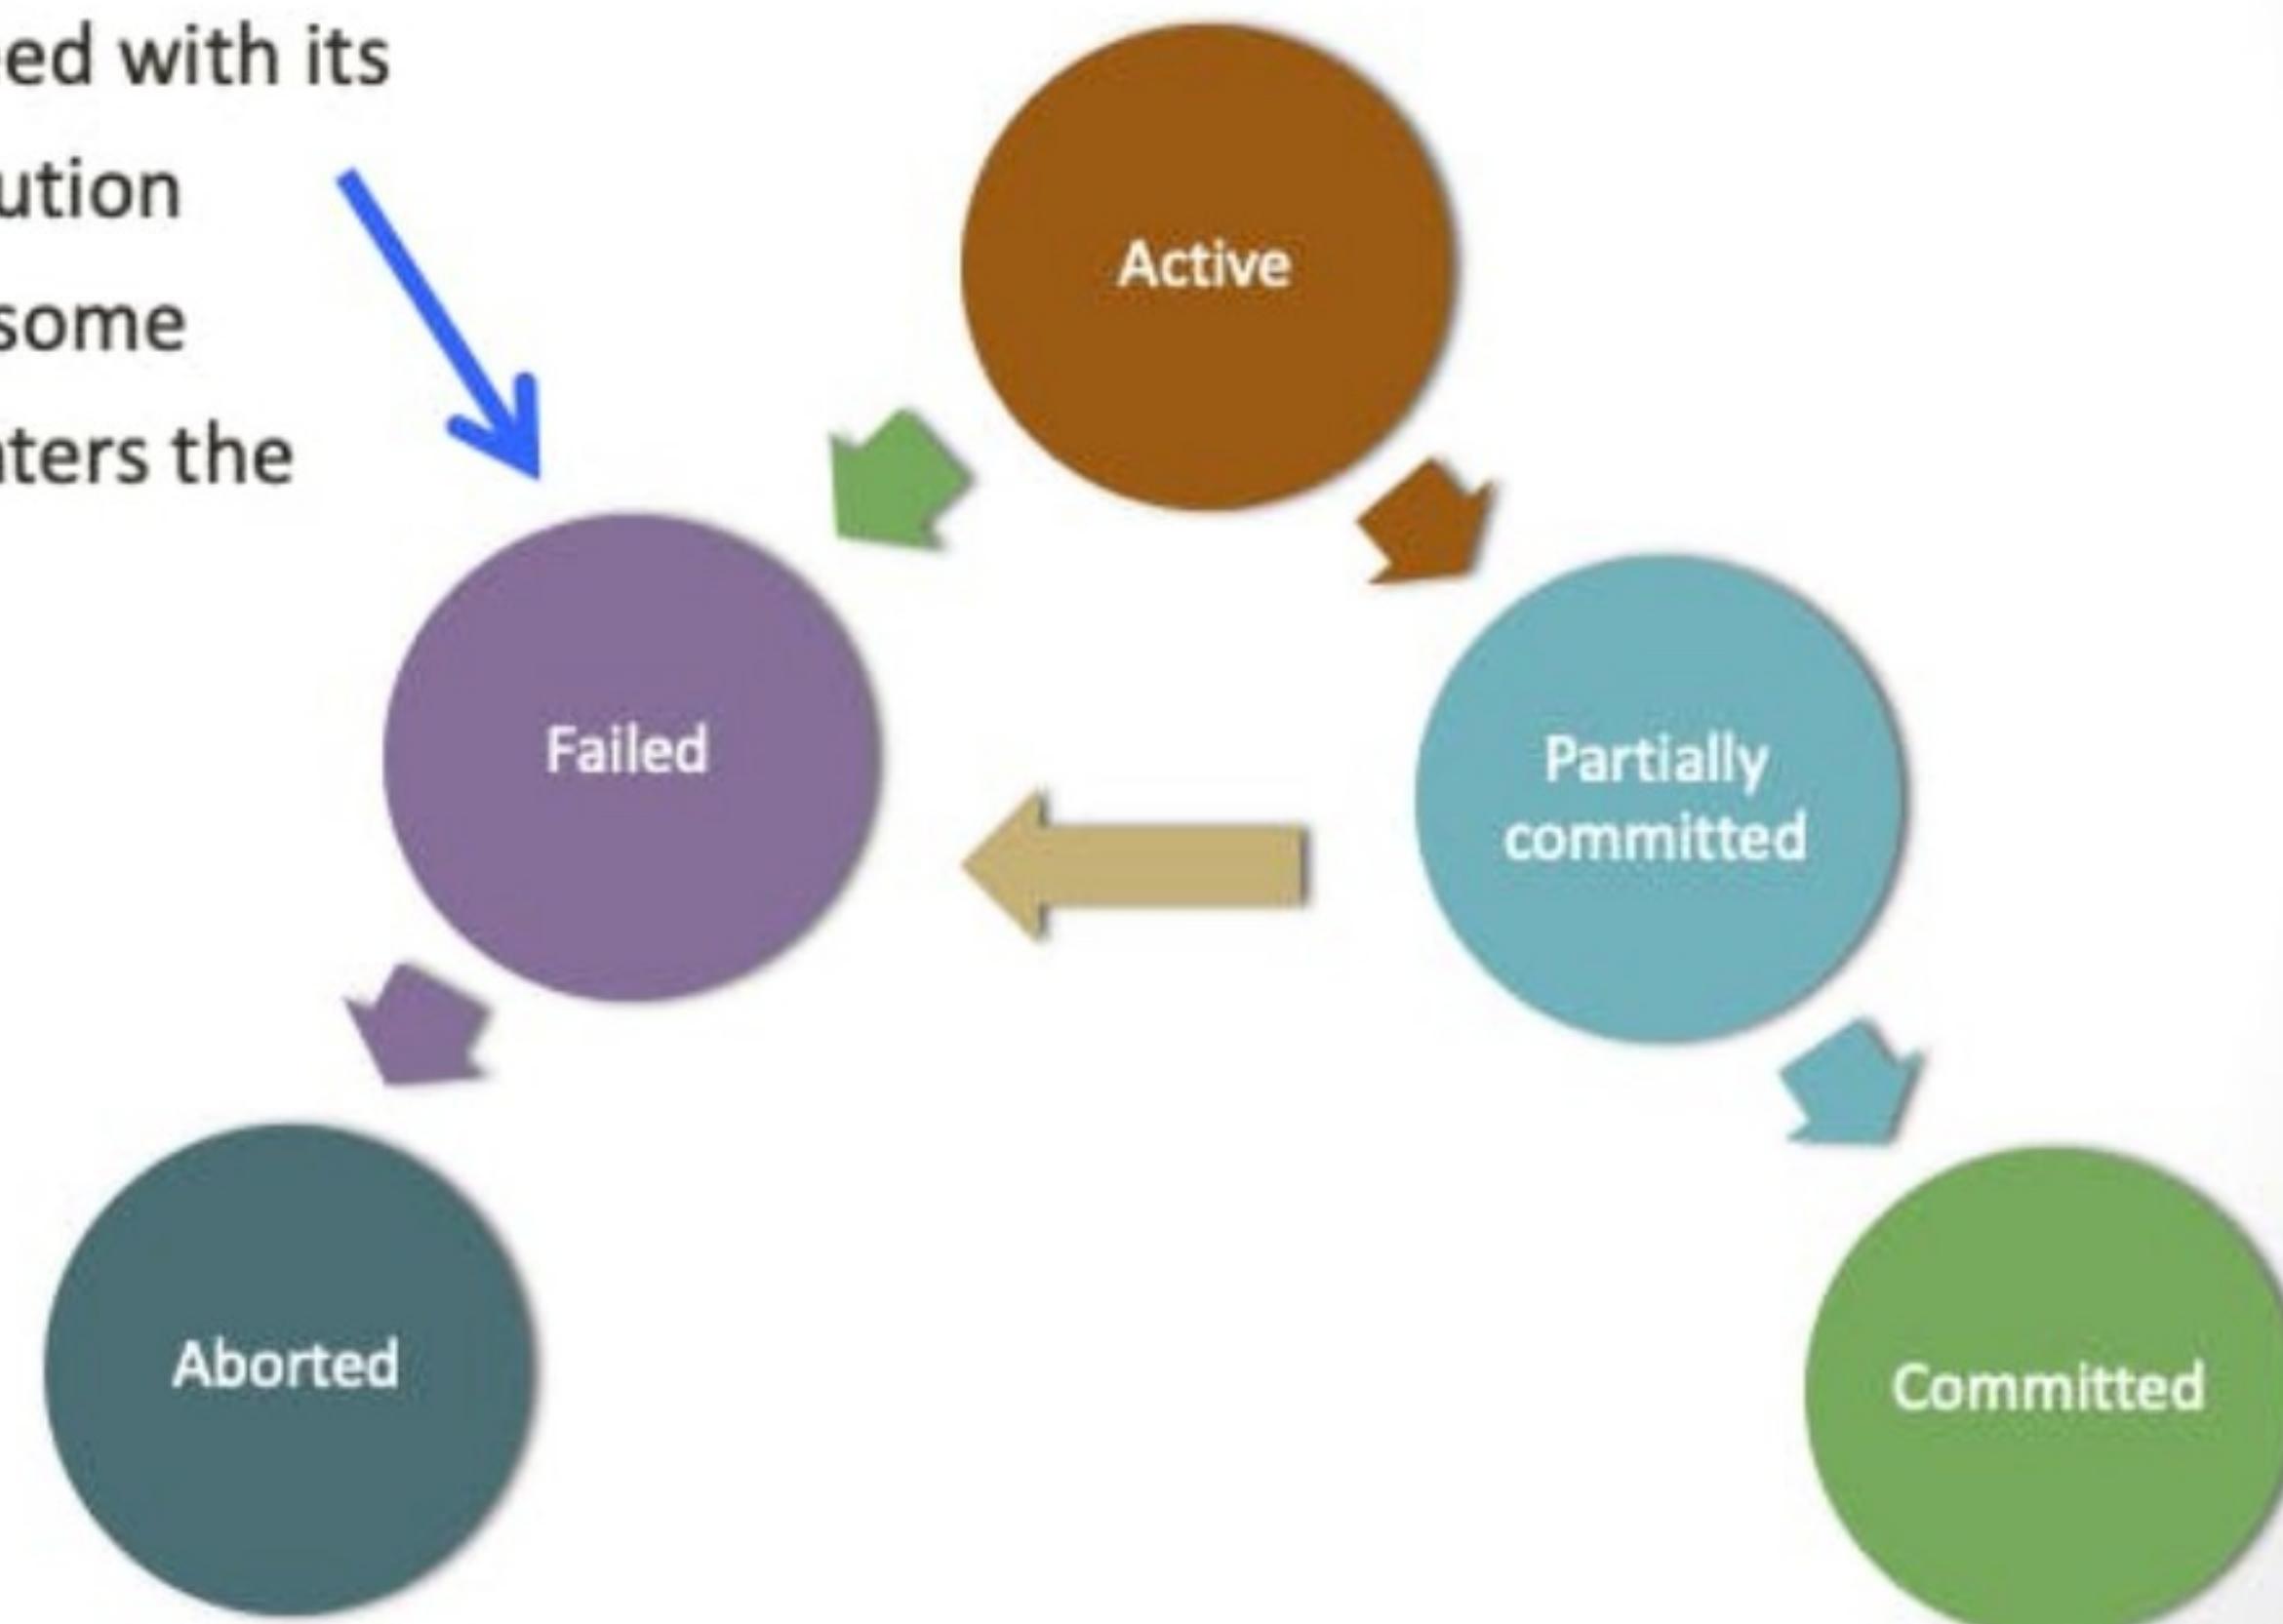
Time	$T_1$	$bal_x$
$t_1$	begin _transaction	100
$t_2$	read( $bal_x$ )	100
$t_3$	$bal_x = bal_x - 10$	100
$t_4$	write( $bal_x$ )	90
$t_5$	end _transaction	90
$t_6$	commit	90

Committed



# Transaction States

- If a transaction can no longer proceed with its normal execution (because of some failure), it enters the **Failed** state.



[ 36 ]



# Transaction States: Example

$T_4$  : Deposit transaction

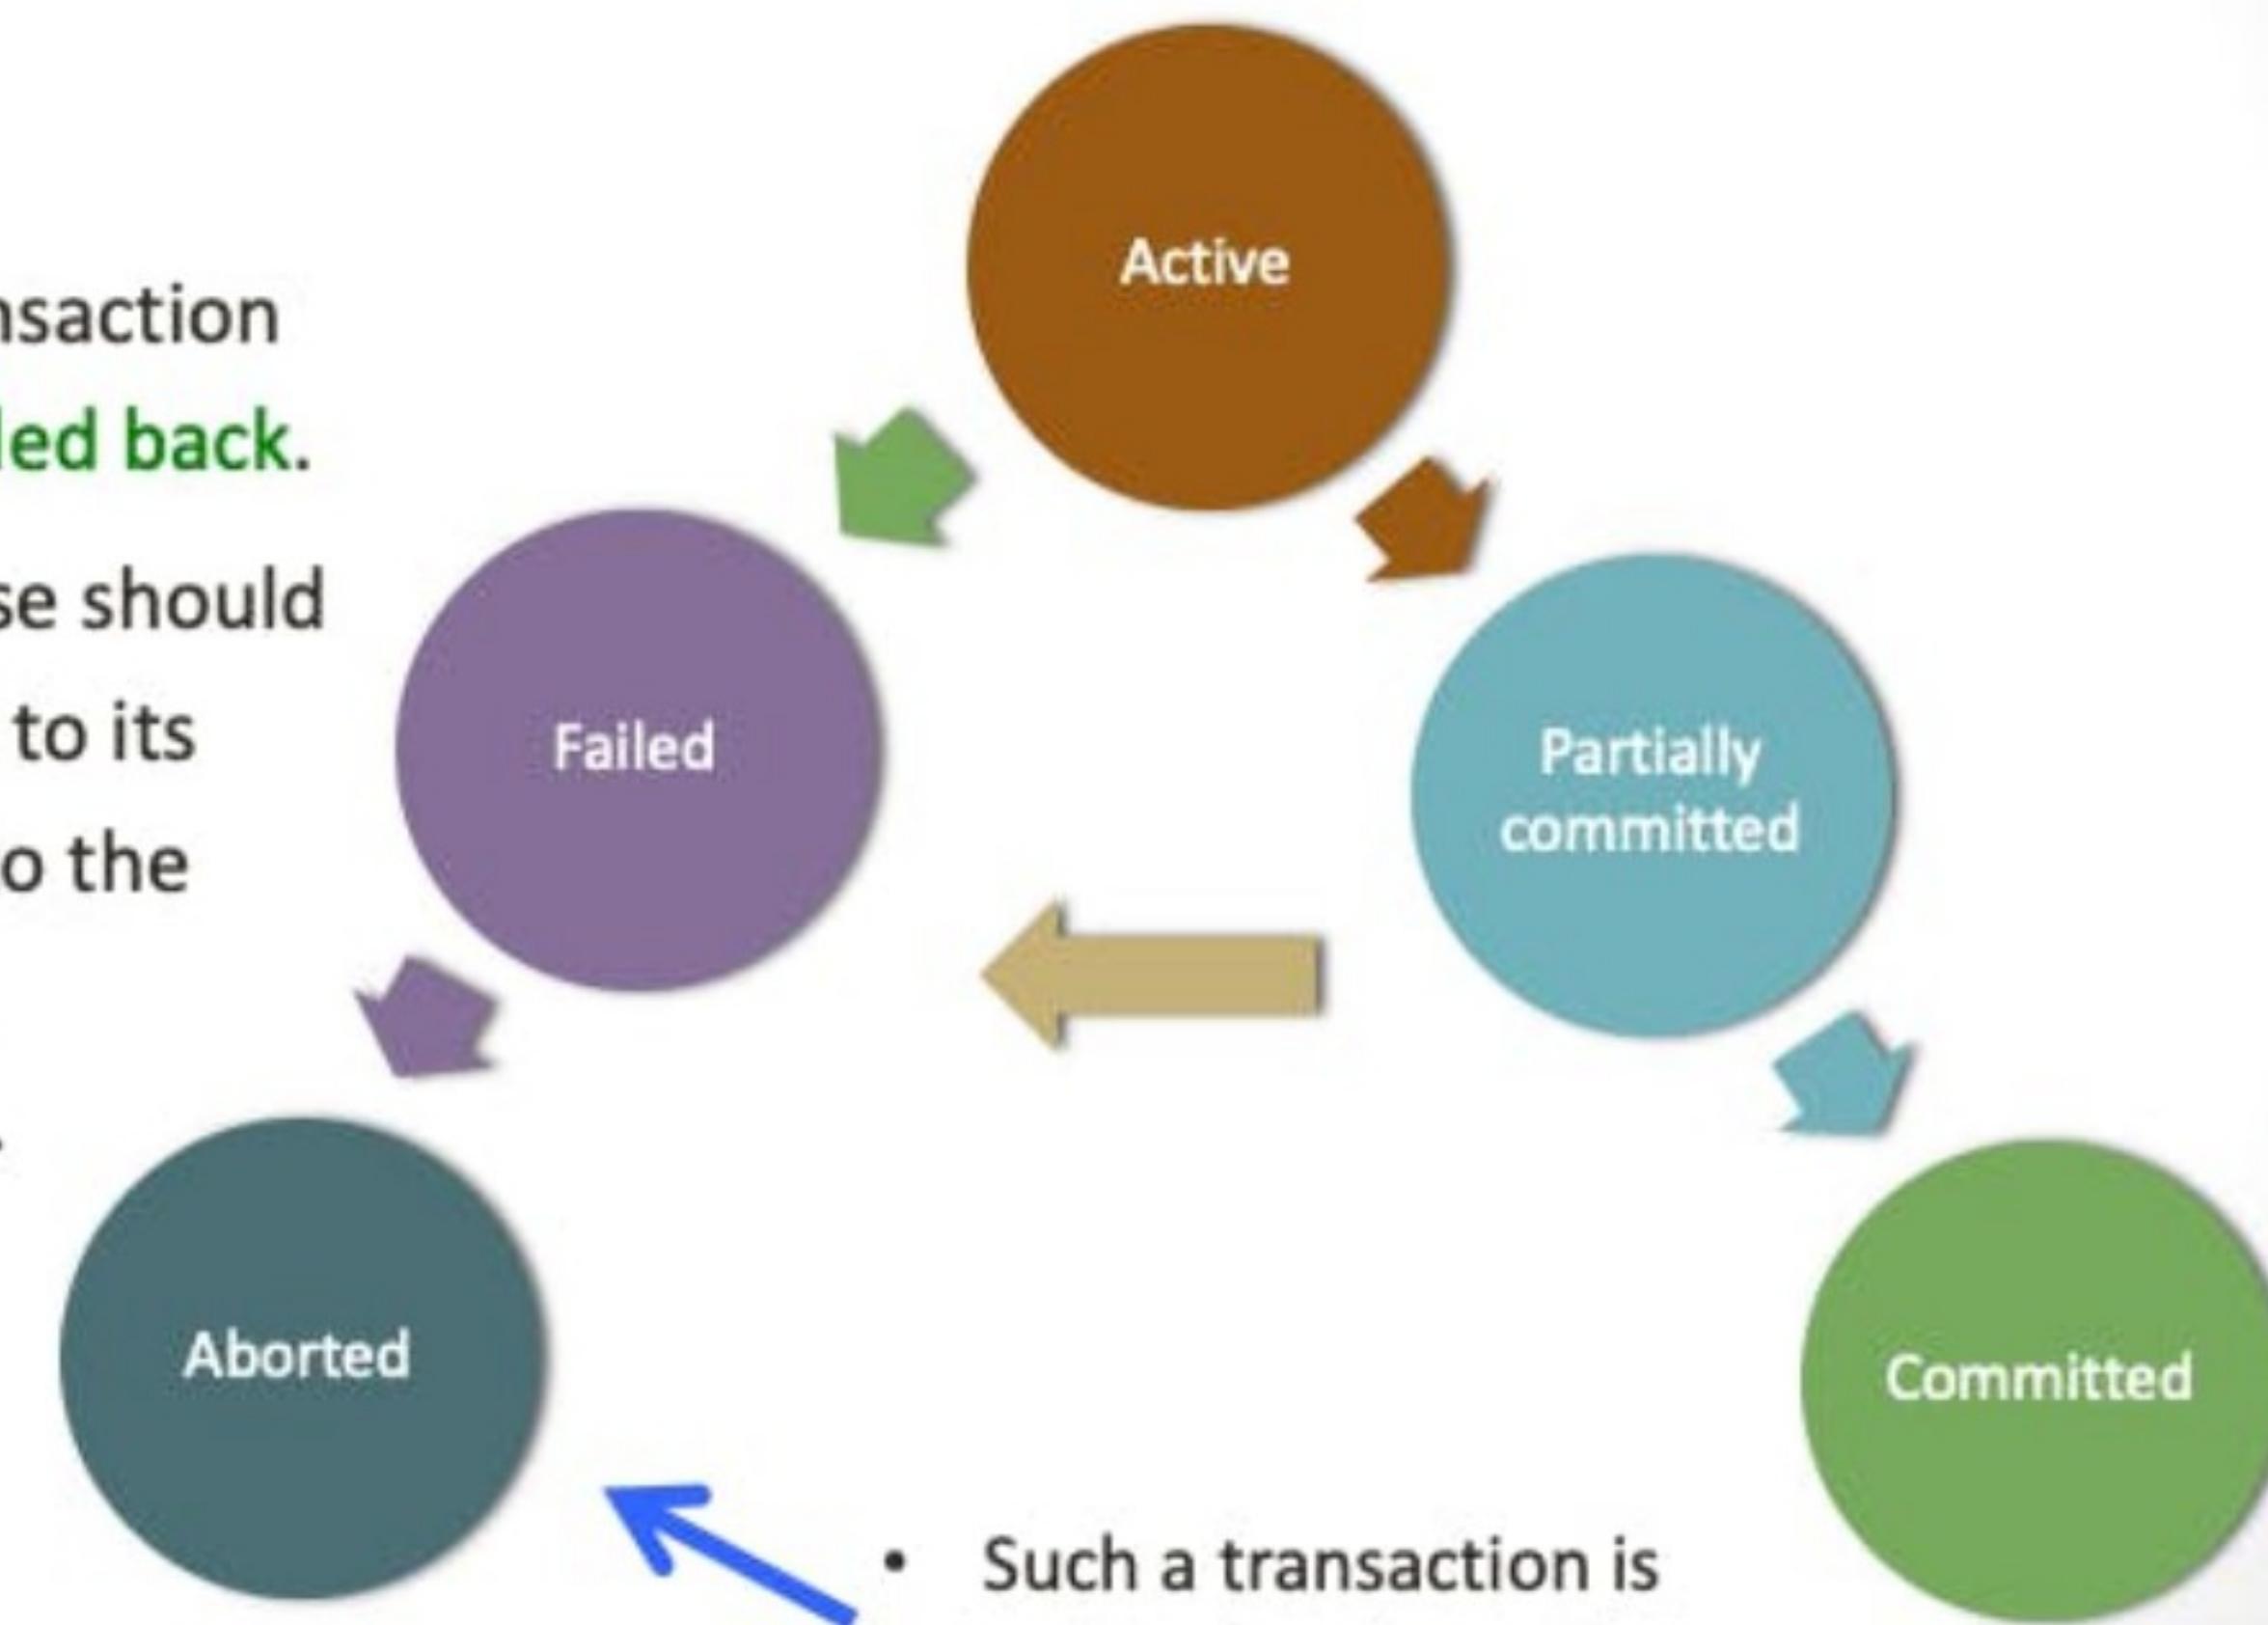
Time	$T_4$	$\text{bal}_x$
$t_1$	begin _transaction	100
$t_2$	read( $\text{bal}_x$ )	100
$t_3$	$\text{bal}_x = \text{bal}_x + 100$	100
$t_4$	write( $\text{bal}_x$ )	200
$t_5$	... (failure)	
$t_6$		
$t_7$		

Failed state



# Transaction States

- A failed transaction must be **rolled back**.
- The database should be restored to its state prior to the start of the transaction.



- Such a transaction is said to be **aborted**.

[ 38 ]



# Transaction States: Example

$T_4$  : Deposit transaction

Time	$T_4$	$\text{bal}_x$
$t_1$	begin _transaction	100
$t_2$	read( $\text{bal}_x$ )	100
$t_3$	$\text{bal}_x = \text{bal}_x + 100$	100
$t_4$	write( $\text{bal}_x$ )	200
$t_5$	... (failure)	
$t_6$	rollback	
$t_7$		



# Transaction Outcomes

- **Committed**
  - The transaction is successfully **completed**.
  - The database reaches a new **consistent state**.
- **Aborted**
  - The transaction does not execute successfully.
  - The **database must be restored to the consistent state it was in** before the transaction started.
  - This is called **rolled back** or **undone**.

[ 40 ]



# Transaction: Other Processes

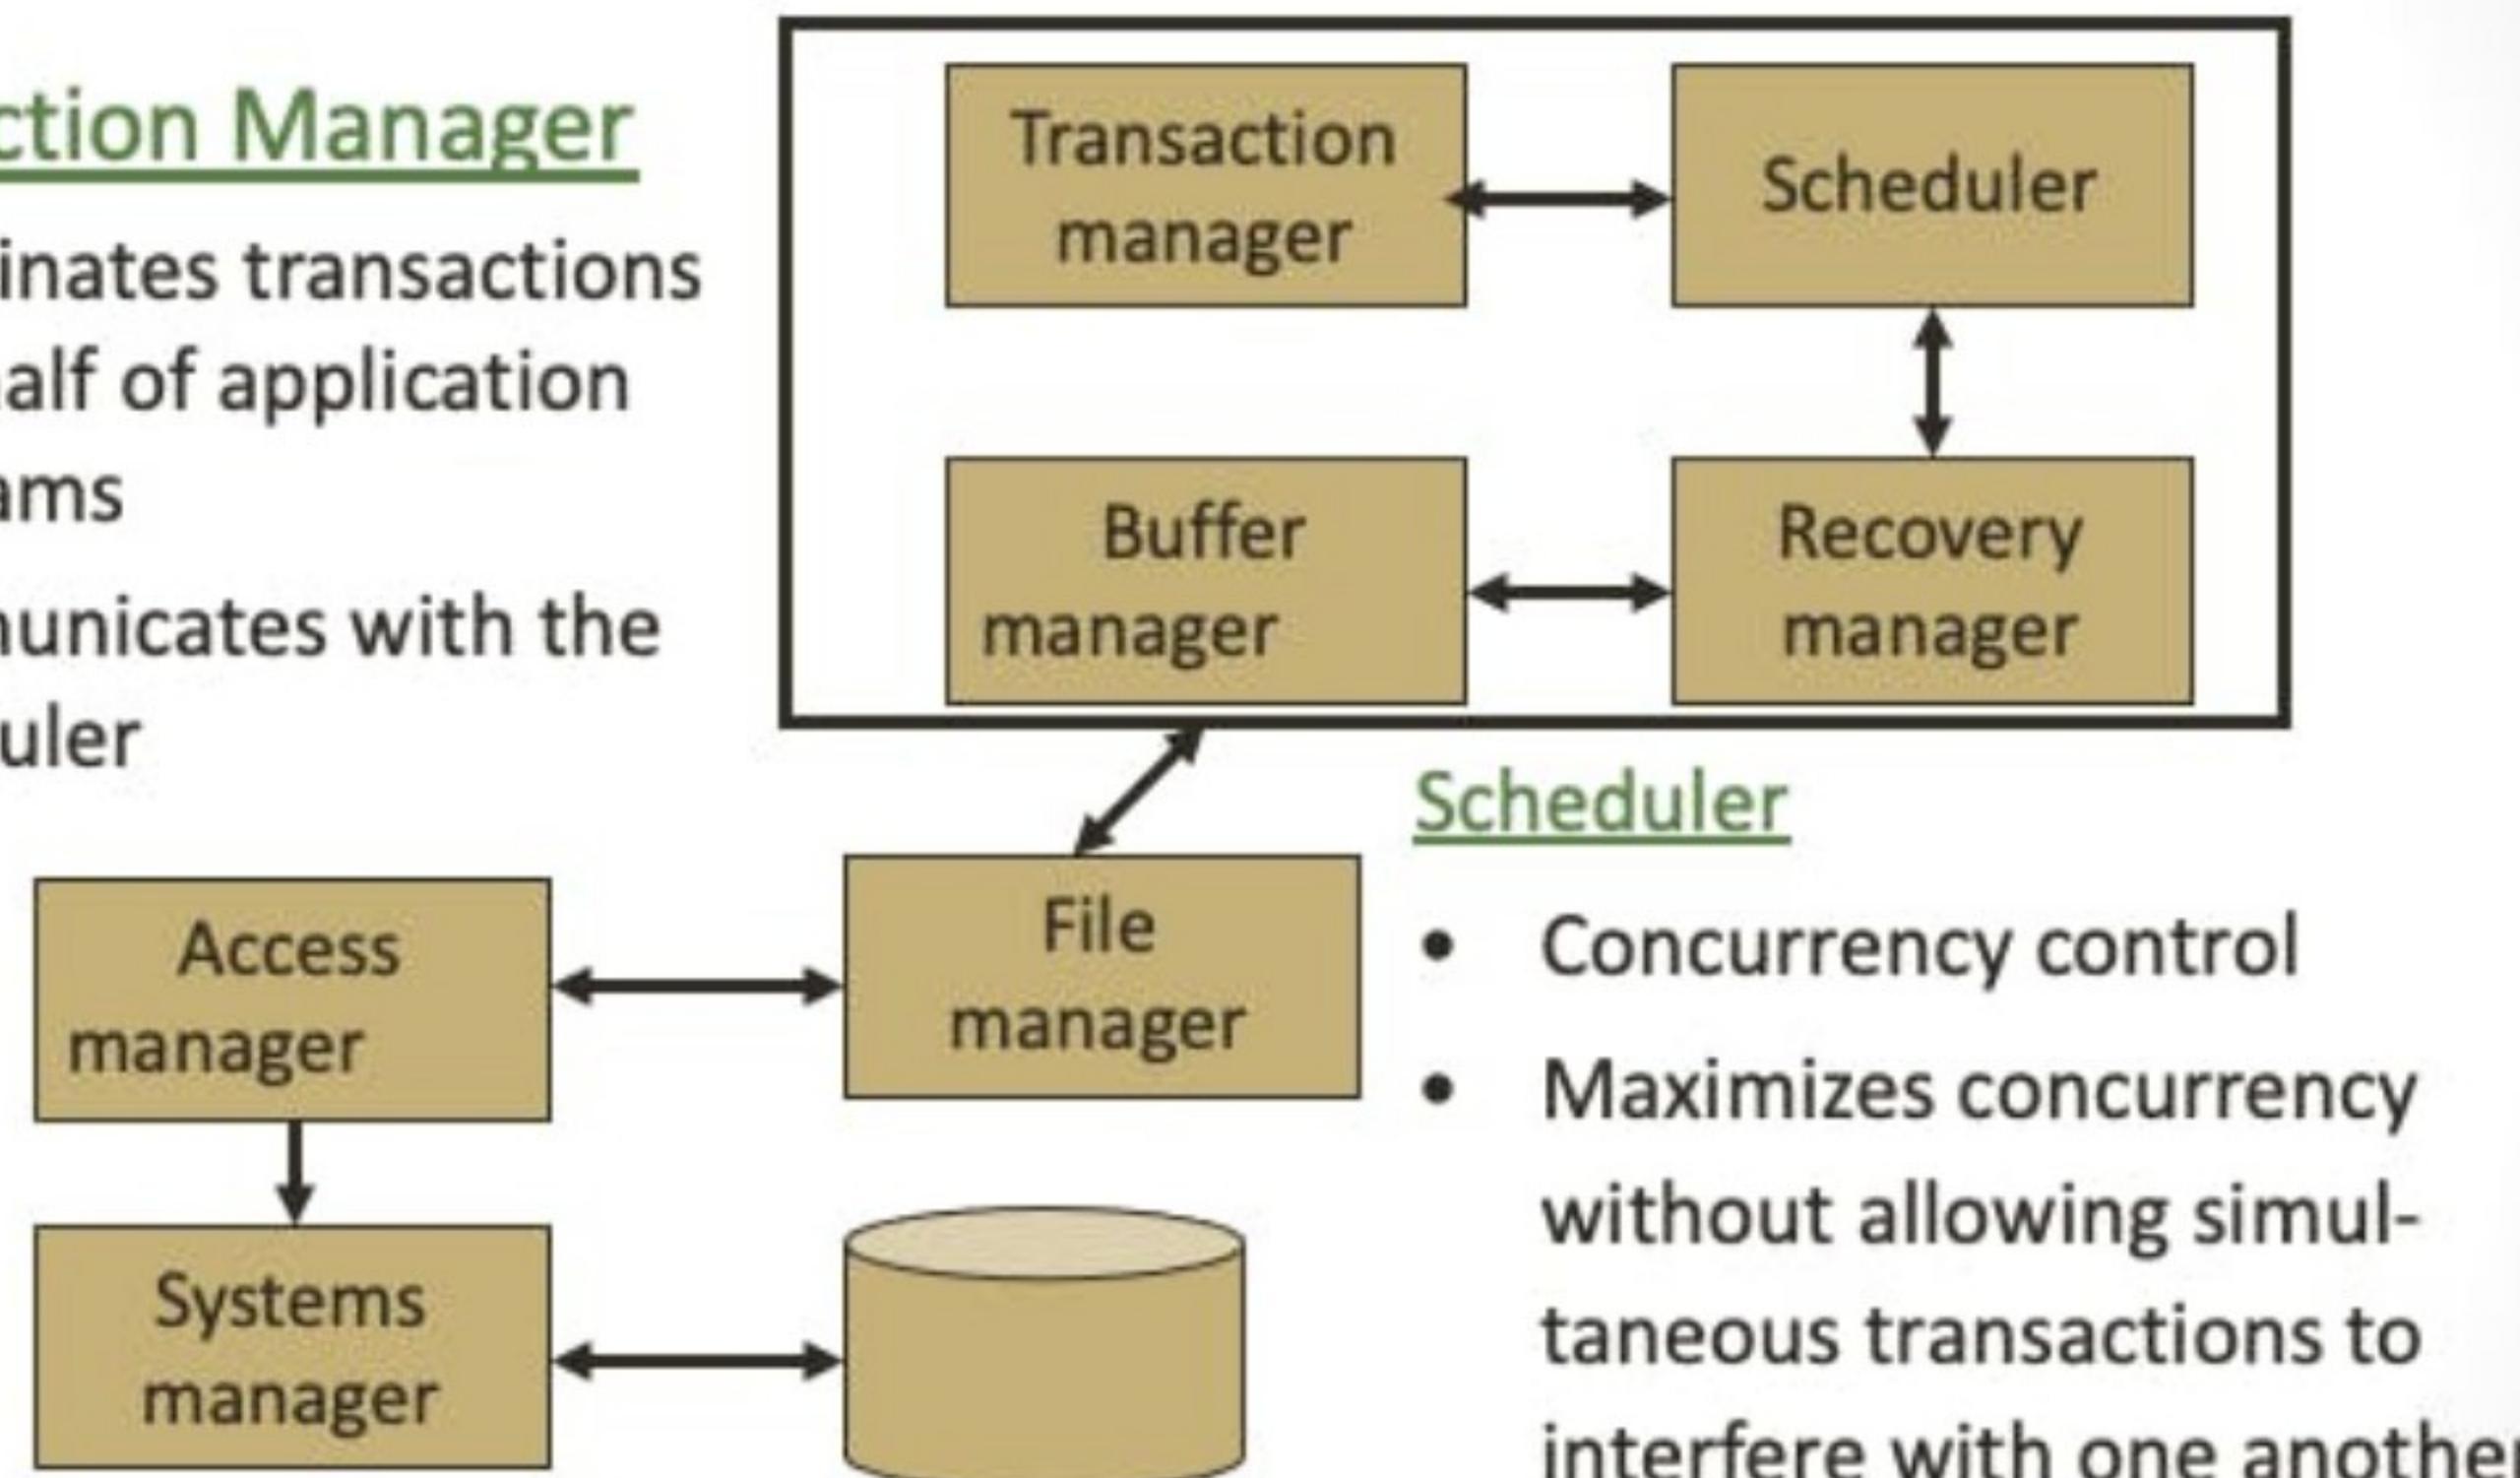
- A committed transaction **cannot** be aborted.
- To reverse the effects of a committed transaction, a **compensating transaction** must be executed.
- An aborted transaction that is rolled back:
  - Can be restarted later
  - Depending on the cause of failure, it may successfully execute and then commit at that time



# DBMS Transaction Subsystem

## Transaction Manager

- Coordinates transactions in behalf of application programs
- Communicates with the Scheduler



## Scheduler

- Concurrency control
- Maximizes concurrency without allowing simultaneous transactions to interfere with one another

- Recovery Manager - ensures that the database is restored to the right state it was in before a failure occurred
- Buffer Manager – handles data transfer between disk and main memory

42

# Levels of Isolation Levels

- Serializable
  - Default
- Repeatable read
  - Only committed records can be read
  - Repeated reads of the same record must return the same value
- Read committed
  - Only committed records can be read
  - Successive reads of record may return different (but committed) values
- Read uncommitted
  - Even uncommitted records may be read

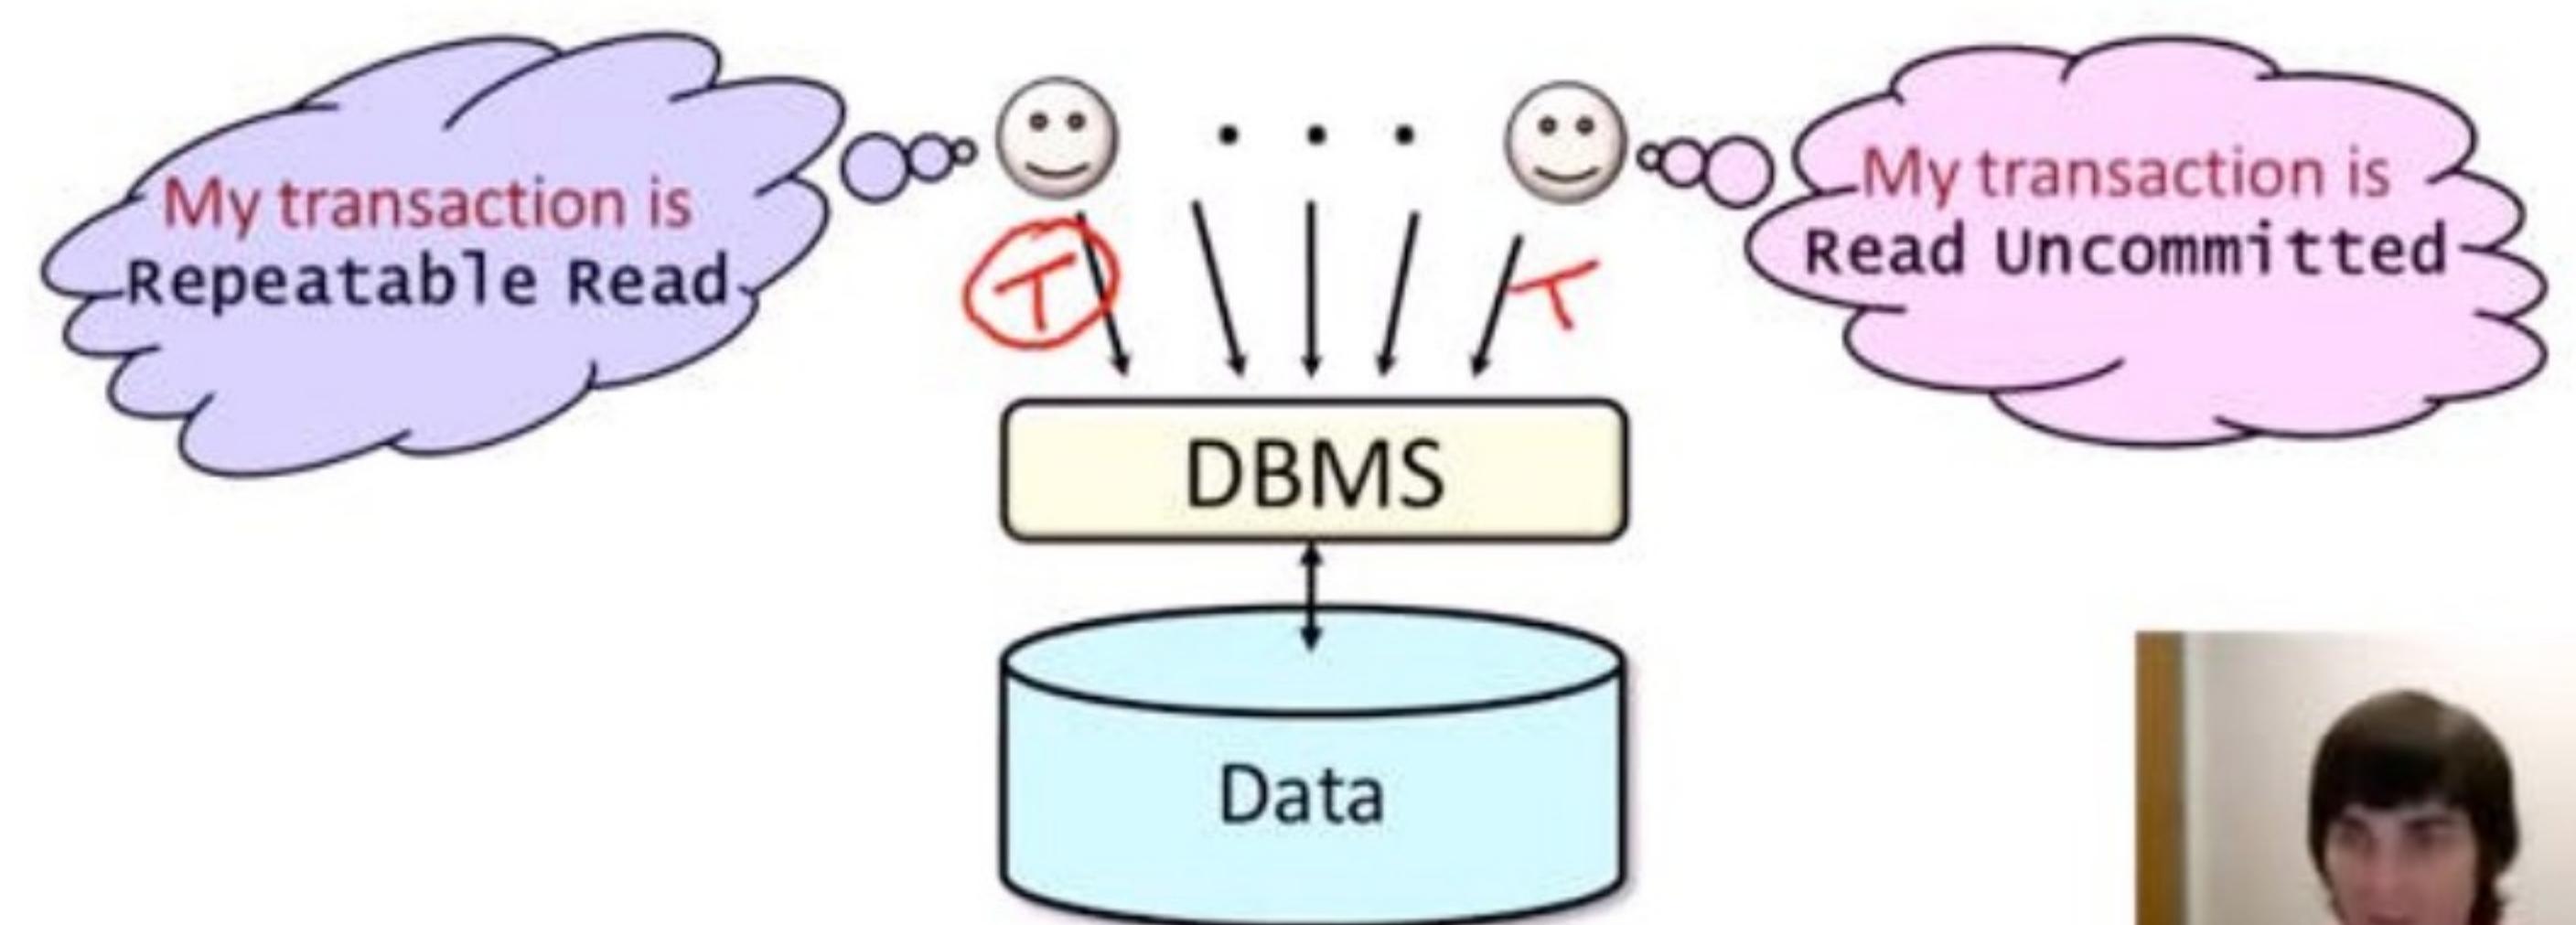
**Decreasing level of DB consistency**

[ 43 ]



# Isolation Levels

- Per transaction
  - Each client can set different isolation levels for its transaction
  - “In the eye of the beholder”



[ 44 ]



# Serializable Schedule

- In a Serializable schedule
  - The execution of interleaved transactions, e.g., T1 and T2, should lead to the same results as executing the transactions serially, either  $\langle T_1, T_2 \rangle$  or  $\langle T_2, T_1 \rangle$

T<sub>1</sub> update Student Set GPA = (1.1) \* GPA where sizeHS > 2500  
concurrent with ...  
T<sub>2</sub> Select Avg(GPA) From Student

$T_1; T_2 \leftarrow$   
 $T_2; T_1$

[ 45 ]



# Dirty Reads

- Dirty data item
  - Written by an uncommitted transaction

↖ 1 `update College Set enrollment=enrollment+1000  
where cName='Stanford'`  
concurrent with ... Commit

↖ 2 `select Avg(enrollment) From College`  
Commit

**Dirty read cannot happen within the same transaction**

↖ 1 `Update Student Set GPA=(1.1)*GPA where sizeHS > 2500`  
concurrent with ...

↖ 2 `Select GPA From Student where SID=123`  
concurrent with ...

↖ 3 `Update Student Set sizeHS = 2600 where SID = 234`

[ 46 ]



# Read Uncommitted

- A transaction may perform dirty reads

```
Update Student Set GPA= (1.1)*GPA where sizeHS > 2500
```

concurrent with ...

```
Set Transaction Isolation Level Read Uncommitted;
Select Avg(GPA) From Student;
```

- In the given example, the Isolation Level setting allows the Select statement to read dirty data if the Update has not committed yet



# Read Committed

## Isolation Level Read Committed

### Transactions

- A transaction may not perform dirty reads  
Still does not guarantee global serializability

T<sub>1</sub> `Update Student Set GPA = (1.1) * GPA Where sizeHS > 2500`  
concurrent with ...

T<sub>2</sub> `Set Transaction Isolation Level Read Committed;`  
`Select Avg(GPA) From Student;`  
`Select Max(GPA) From Student;`

- In the given example, the 2 SELECT statements may read committed but different GPA values if the reads were performed before and after the Update statement.
- Is there an equivalent serial schedule?

[ 48 ]

# Repeatable Read

- A transaction may not perform dirty reads
- An item read multiple times cannot change value
- Still does not guarantee global serializability

T1 Update Student Set GPA = (1.1) \* GPA;  
Update Student Set sizeHS = 1500 where SID = 123;

concurrent with ...

T2 Set Transaction Isolation Level Repeatable Read;  
Select Avg(GPA) From Student;  
Select Avg(sizeHS) From Student;

- Why is the schedule still not serializable despite not performing dirty reads?



# Repeatable Read

- A relation can change while being read multiple times through “phantom” tuples

```
T1: Insert Into Student [ 100 new tuples ]
concurrent with ...
T2: Set Transaction Isolation Level Repeatable Read;
      Select Avg(GPA) From Student;
      Select Max(GPA) From Student;
```

- Repeatable Read allows the AVG operation in T2 to be executed before T1 and MAX to be executed after T1.
- Thus, if the second SELECT statement in T2 is another AVG operation, then the results would not be the same.
- Phantom tuples are those that emerge during an execution.

[ 50 ]



# Repeatable Read

- The previous case where 2 separate Read operations of T2 were allowed to be executed before and after T1 can only happen during Insert, but not Delete. Why?

A diagram illustrating a database transaction scenario. At the top, a box contains the SQL command: `Delete From student [ 100 tuples ]`. Below it, the text "concurrent with . ." is followed by another box containing:  
`Set Transaction Isolation Level Repeatable Read;`  
`Select Avg(GPA) From Student;`  
`Select Max(GPA) From Student;`

Annotations with pink arrows and circles highlight the "Delete" command, the "100 tuples" count, the "concurrent with" note, and the "GPA" columns in the select statements.

[ 51 ]



# Isolation Levels: Summary

- Weaker Isolation Levels
  - Increased concurrency + Decreased overhead = Increased performance
  - Weaker consistency guarantees
  - Some systems have default **Repeatable Read** (e.g., Oracle, MySQL)
- Each transaction's reads must conform to its Isolation level

*weak*

	dirty reads	non-repeatable reads	phantoms
Read Uncommitted	Y	Y	Y
Read Committed	N	Y	Y
Repeatable Read	N	N	Y
Serializable	N	N	N

[ 52 ]



# Transaction Definition in SQL

- SQL standard
  - A transaction begins implicitly
- To end a transaction
  - Commit [work] – commits the current transaction and begins a new one
  - Rollback [work] – causes the current transaction to abort
- By default in almost all DBMS, every SQL statement commits implicitly if it executes successfully
  - To turn off Implicit Commit, e.g., in JDBC

```
connection.setAutoCommit(false);
```

[ 53 ]



# References

- Read the following:

**Chapter 14: Transactions**

**Chapter 15: Concurrency**

**Chapter 16: Recovery**

Silberschatz, A., Korth, H. & Sudarshan, S. (2010). *Database System Concepts*, 6th Edition. McGraw-Hill Book Co.

**Chapter 22: Transaction Management**

Connolly, T. & Begg, C. (2015). *Database Systems: A Practical Approach to Design, Implementation, and Management*, 6th Edition. Harlow, Essex: Addison-Wesley

- Watch the videos of Jennifer Widom (Stanford University), eg,

**Introduction to Transactions**

**Transactions – Isolation Levels**

[ 54 ]

